

# GraphOps: A Dataflow Library for Graph Analytics Acceleration

Tayo Oguntebi\*  
Google, Inc.

Kunle Olukotun  
Pervasive Parallelism Laboratory  
Stanford University

\* Work done while the author  
was at Stanford University

22 February 2016  
FPGA 2016

# Outline

- ❖ The GraphOps Library
- ❖ Locality-Optimized Graph Representation
- ❖ Results and Conclusions

# The GraphOps Library

Betweenness Centrality

```
// BFS order iteration from s  
InBFS(v: G.Nodes From s) {
```

Opt

```
Do {
```

PageRank

Conductance

```
Din = Sum(u:G.Nodes) (u.member == num) {  
  u.Degree()  
};
```

```
Dout = Sum(u:G.Nodes) (u.member != num) {  
  u.Degree()  
};
```

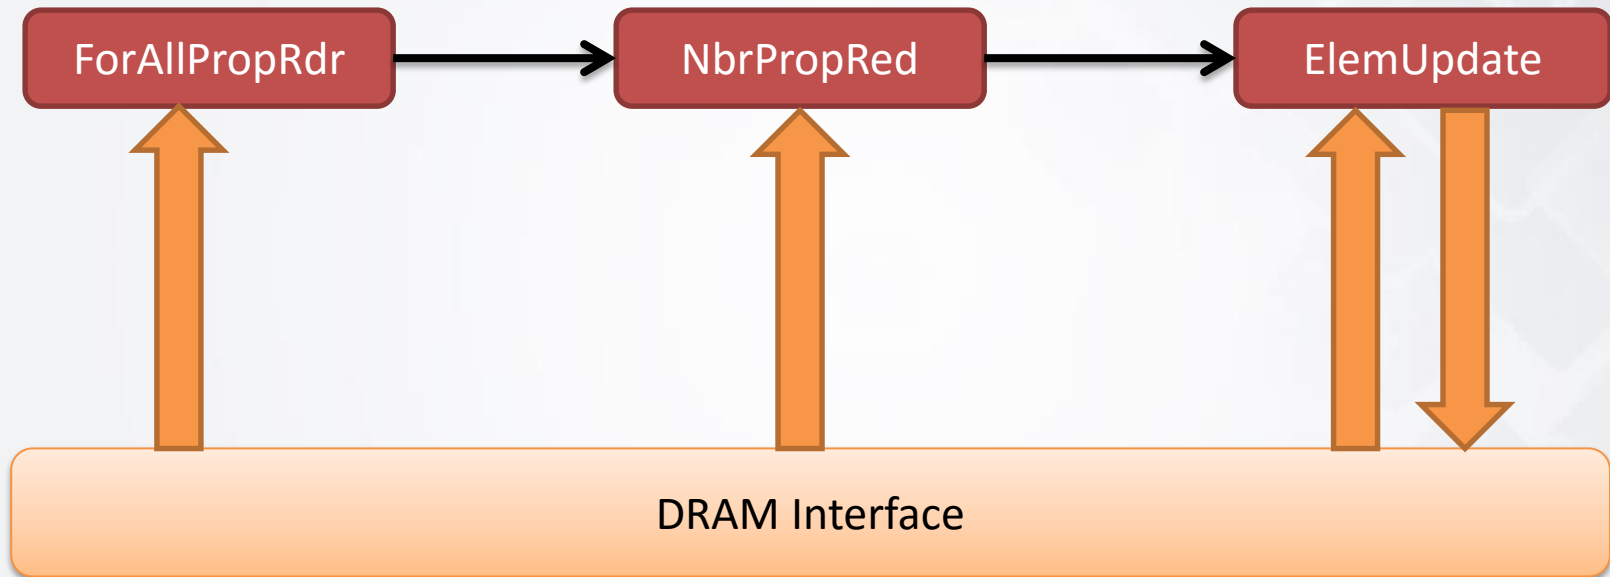
```
Cross = Sum(u:G.Nodes) (u.member == num) {  
  Count(j:u.Nbrs) (j.member != num)  
};
```

to s

# Running Example: PageRank

```
Procedure pagerank()  
{  
    Double diff;  
    Int cnt = 0;  
    Double N = G.NumNodes();  
    G.pg_rank = 1 / N;  
    Do {  
        diff = 0.0;  
        Foreach (t: G.Nodes) {  
            Double val = (1-d) / N + d*  
                Sum(w: t.InNbrs) {  
                    w.pg_rank / w.OutDegree();  
                } ;  
            diff += | val - t.pg_rank |;  
            t.pg_rank <= val @ t;  
        }  
        cnt++;  
    } While ((diff > e) && (cnt < max));  
}
```

# The GraphOps Library



# The GraphOps Library

DATA

CONTROL

UTILITY

ForAllPropRdr

NbrPropRed

ElemUpdate

AllNodePropRdr

NbrPropRdr

SetReader

SetWriter

NbrPropFilter

GlobNbrRed

VertexReader

NbrSetReader

# The GraphOps Library

DATA

CONTROL

UTILITY

## Data Readers

ForAllPropRdr

VertexReader

NbrSetReader

AllNodePropRdr

NbrPropRdr

## Reduction

GlobNbrRed

NbrPropRed

## Set Manipulation

SetReader

SetWriter

## Property Filtering

NbrPropFilter

## Mutation

ElemUpdate

# The GraphOps Library

- Set of optimized hardware blocks for executing common graph processing functions
  - High-level: Easy to use
  - Composable: Flexible enough to compose different applications
  - Extensible and parameterizable
  - Pre-verified: Low-level implementation details built-in to the design

**Problem: Poor Locality → Poor Performance!**





# Outline

- ❖ The GraphOps Library
- ❖ Locality-Optimized Graph Representation
- ❖ Results and Conclusions

# Rethinking the Graph Representation


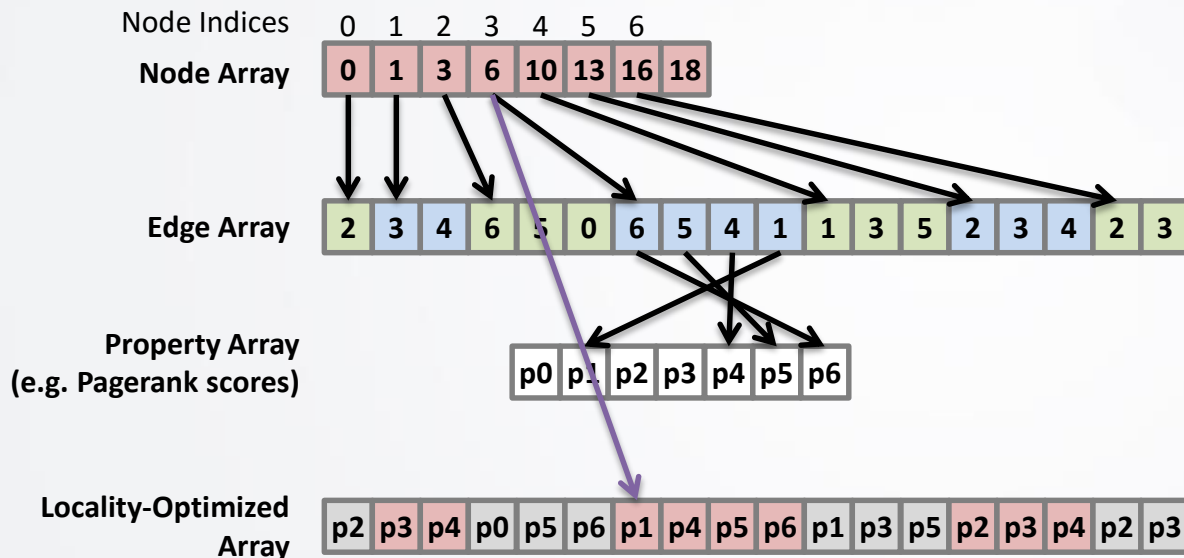
## Conventional Form: Compressed Sparse Row (Adjacency Lists)



**No locality!**

# Rethinking the Graph Representation


## Locality-Optimized Form



Trades off compactness for locality...  
Space for time

# Pre-Processing the Layout

- We have locality...now need to restore consistency



- ProcessGraphLayout(): *Scatter* operation
  - Performed on the host

“The cheapest decent memory controller that you can buy is still an Intel Xeon CPU...” – Prof. Christos Kozyrakis


# Programming Model

```
Graph* g;  
GenerateGraph(g);  
  
PreprocessGraphLayout();    // Prepare locality-optimized form  
  
do {  
    WriteToDeviceMem();  
    Run();  
    ReadFromDeviceMem();  
  
    ProcessGraphLayout();    // i.e. scatter  
} while (not converged);
```

# Outline

- ❖ The GraphOps Library
- ❖ Locality-Optimized Graph Representation
- ❖ Results and Conclusions

# Energy Efficiency (Throughput / Watt)



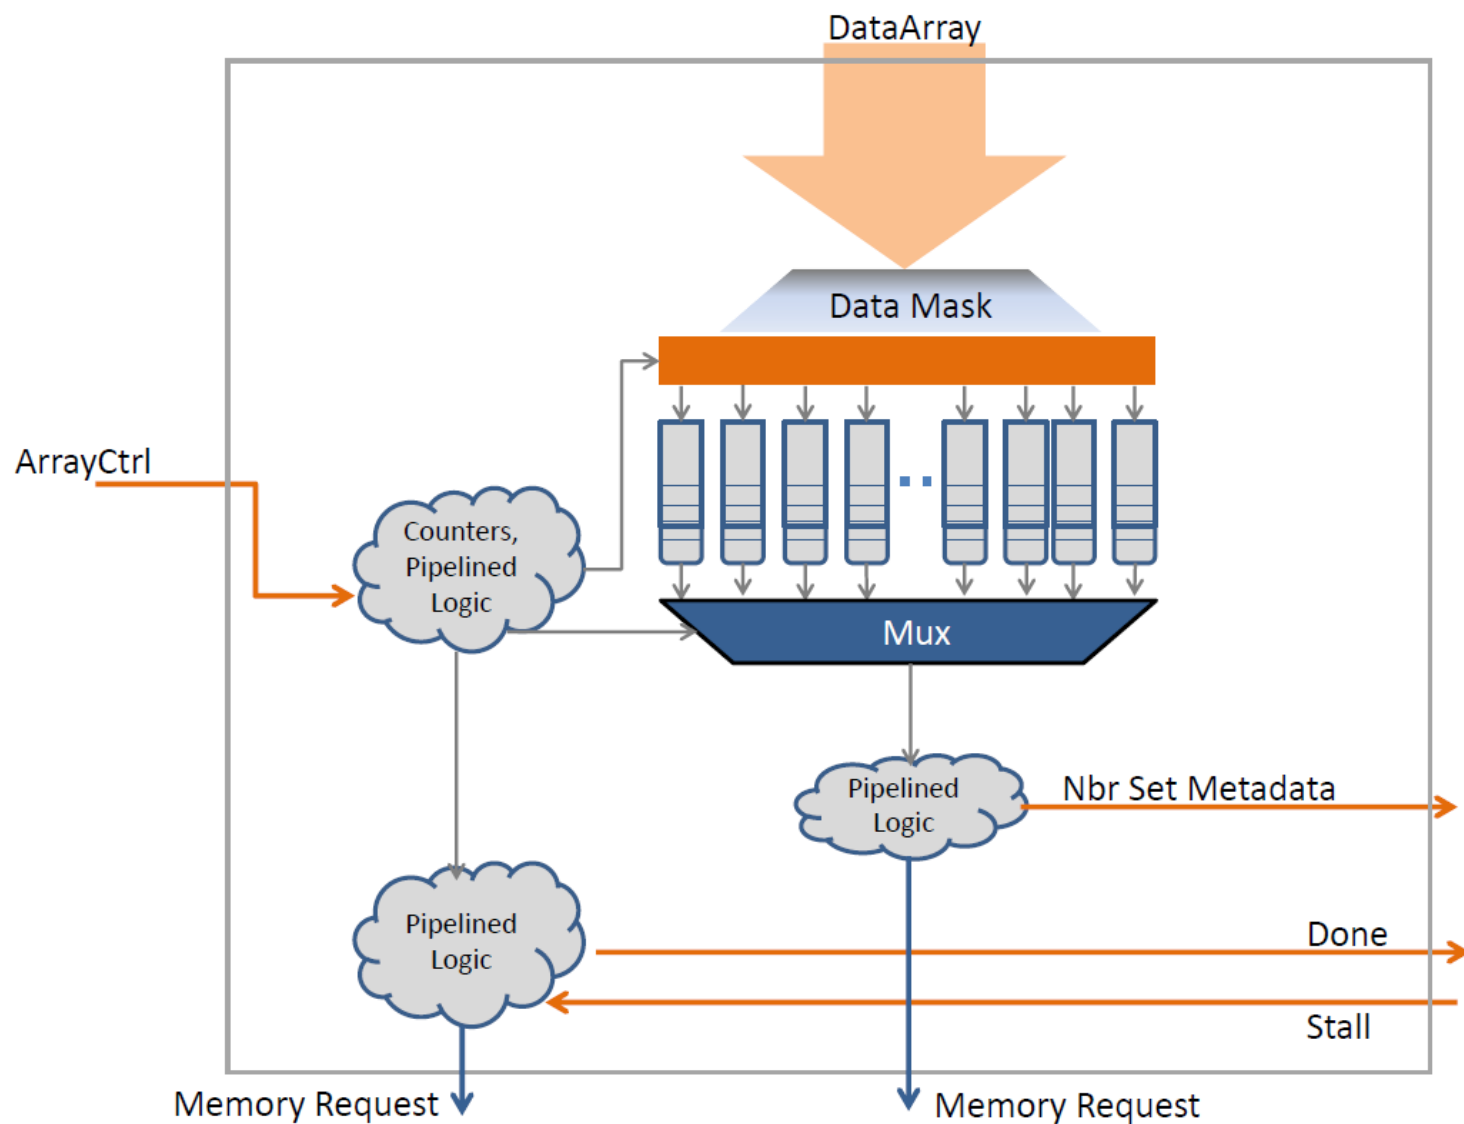
# Thank You

- Details and full results in the paper
- Questions: Find me during the break / poster session.
- Complete library open-sourced (MIT License) and available at:  
<https://github.com/tayo/GraphOps>


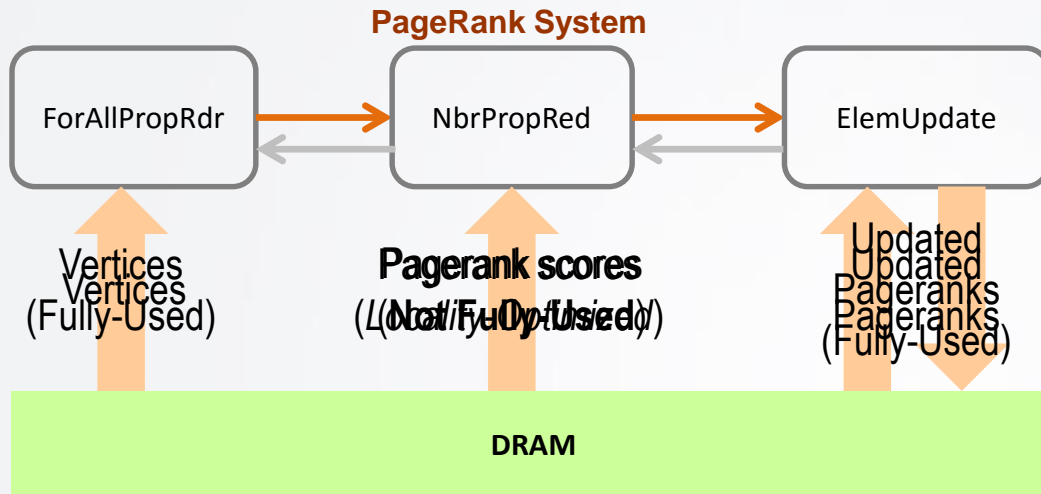


# Supplementary Material

# ForAll Property Reader



# Bandwidth Usage



## Evaluation Platforms

- Intel Xeon 5650 @ 2.7GHz
  - 2 sockets, 12 cores, 24 threads
  - **Bandwidth: 32 GB/s per socket**
  - **3 Memory Channels**
- FPGA: Xilinx Virtex-6 (150MHz)
  - Connected to host via PCIe8 Gen 2
  - **Bandwidth: 38.4 GB/s**

Effective performance is about 1/6 of what bandwidth allows

Single-memory channel

L-O array access has to wait on others

## Constraining Factors

### Locality

- Optimal: Sequential access
- Using: Alternating reads to the different arrays. All units operating simultaneously

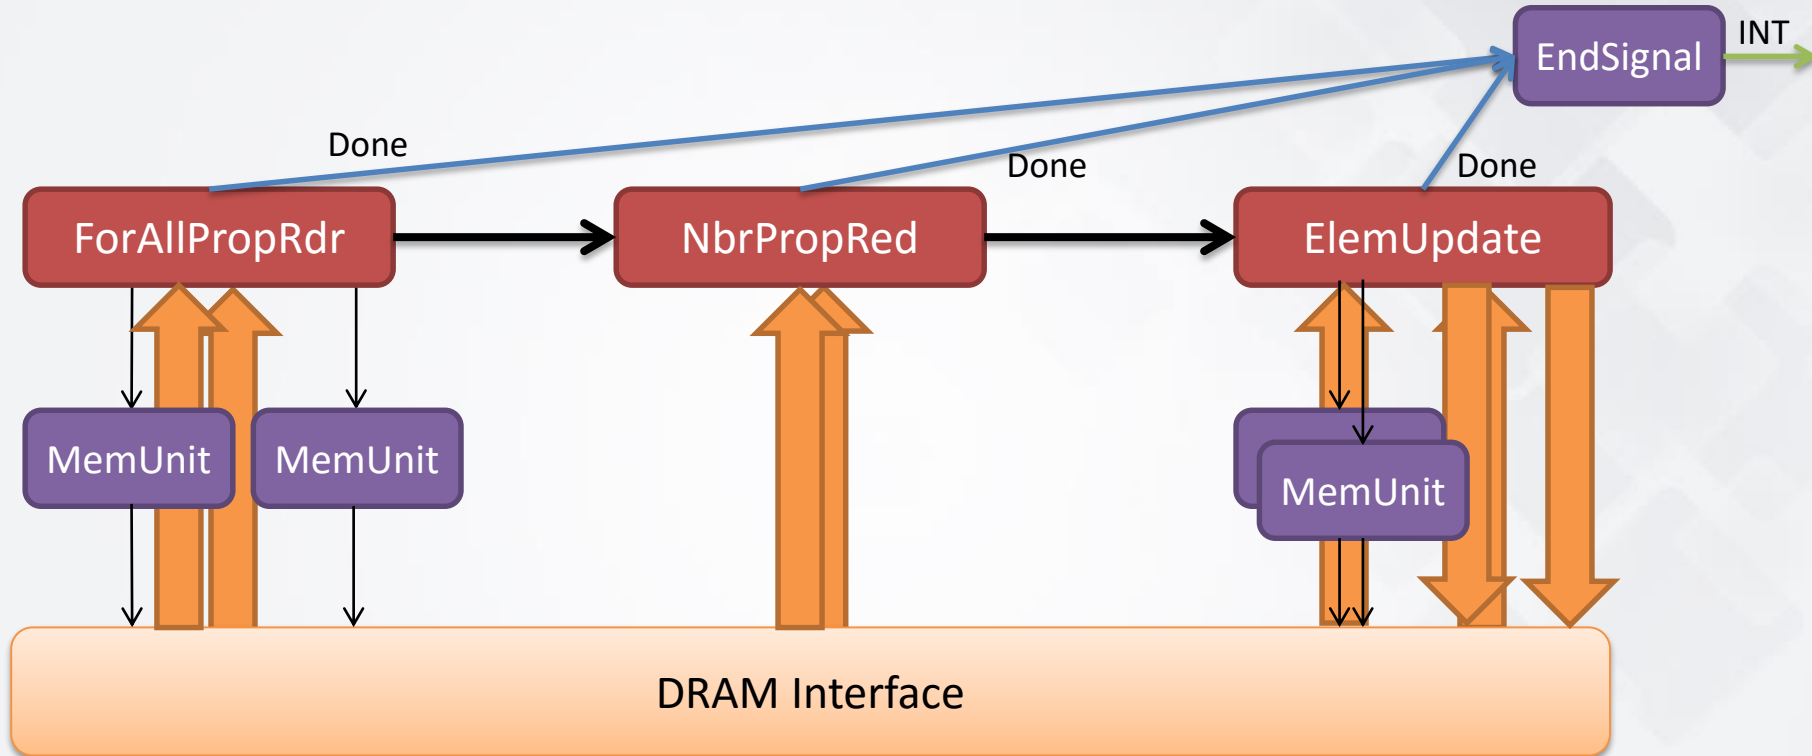
### Packet size

- Optimal: 384 bytes x 4
- Using: 192 bytes x 2

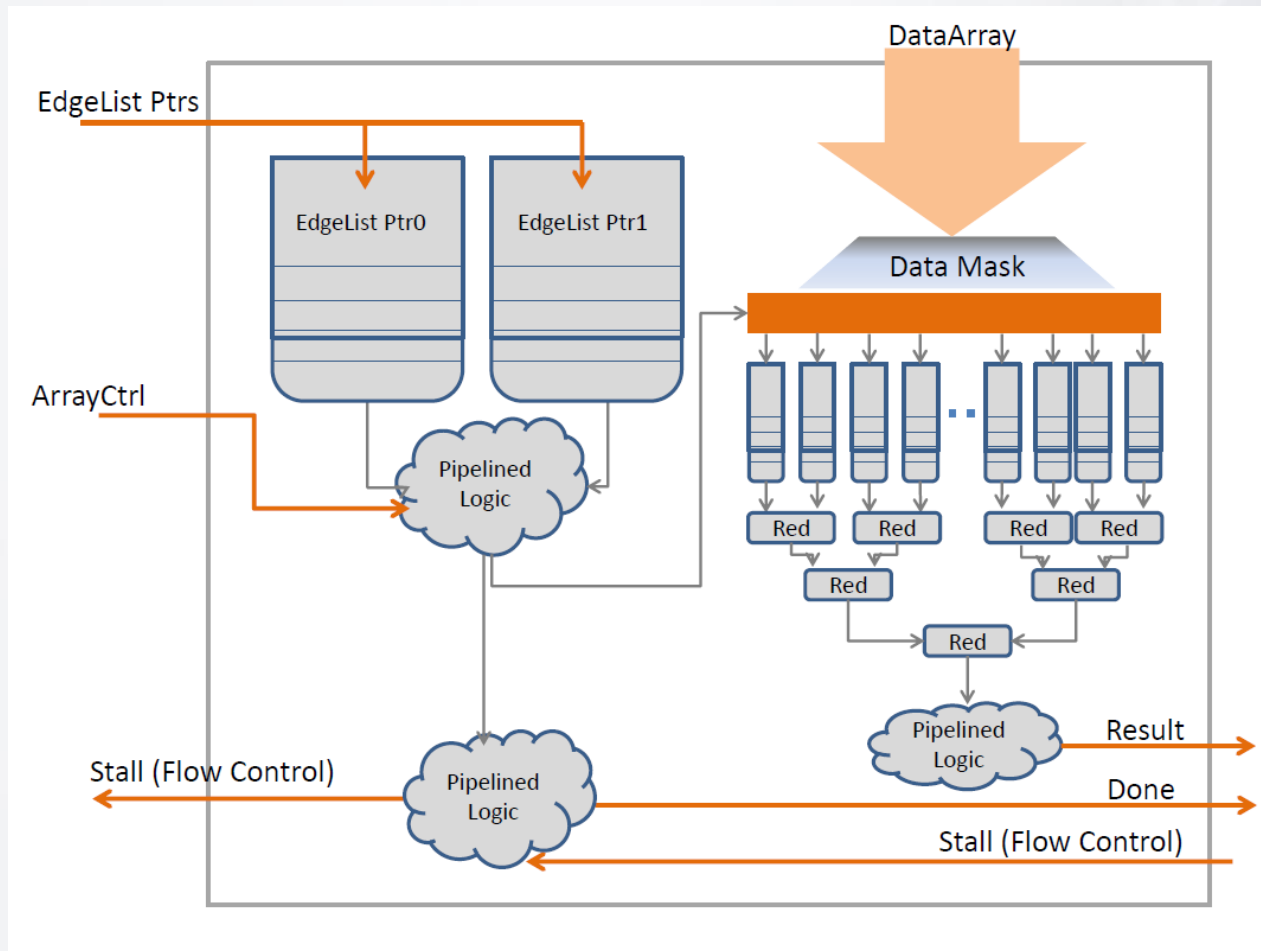
### Burst size

- Optimal: as large as possible (max 256)
- Using: usually 1-2 (enough for a nbr set)

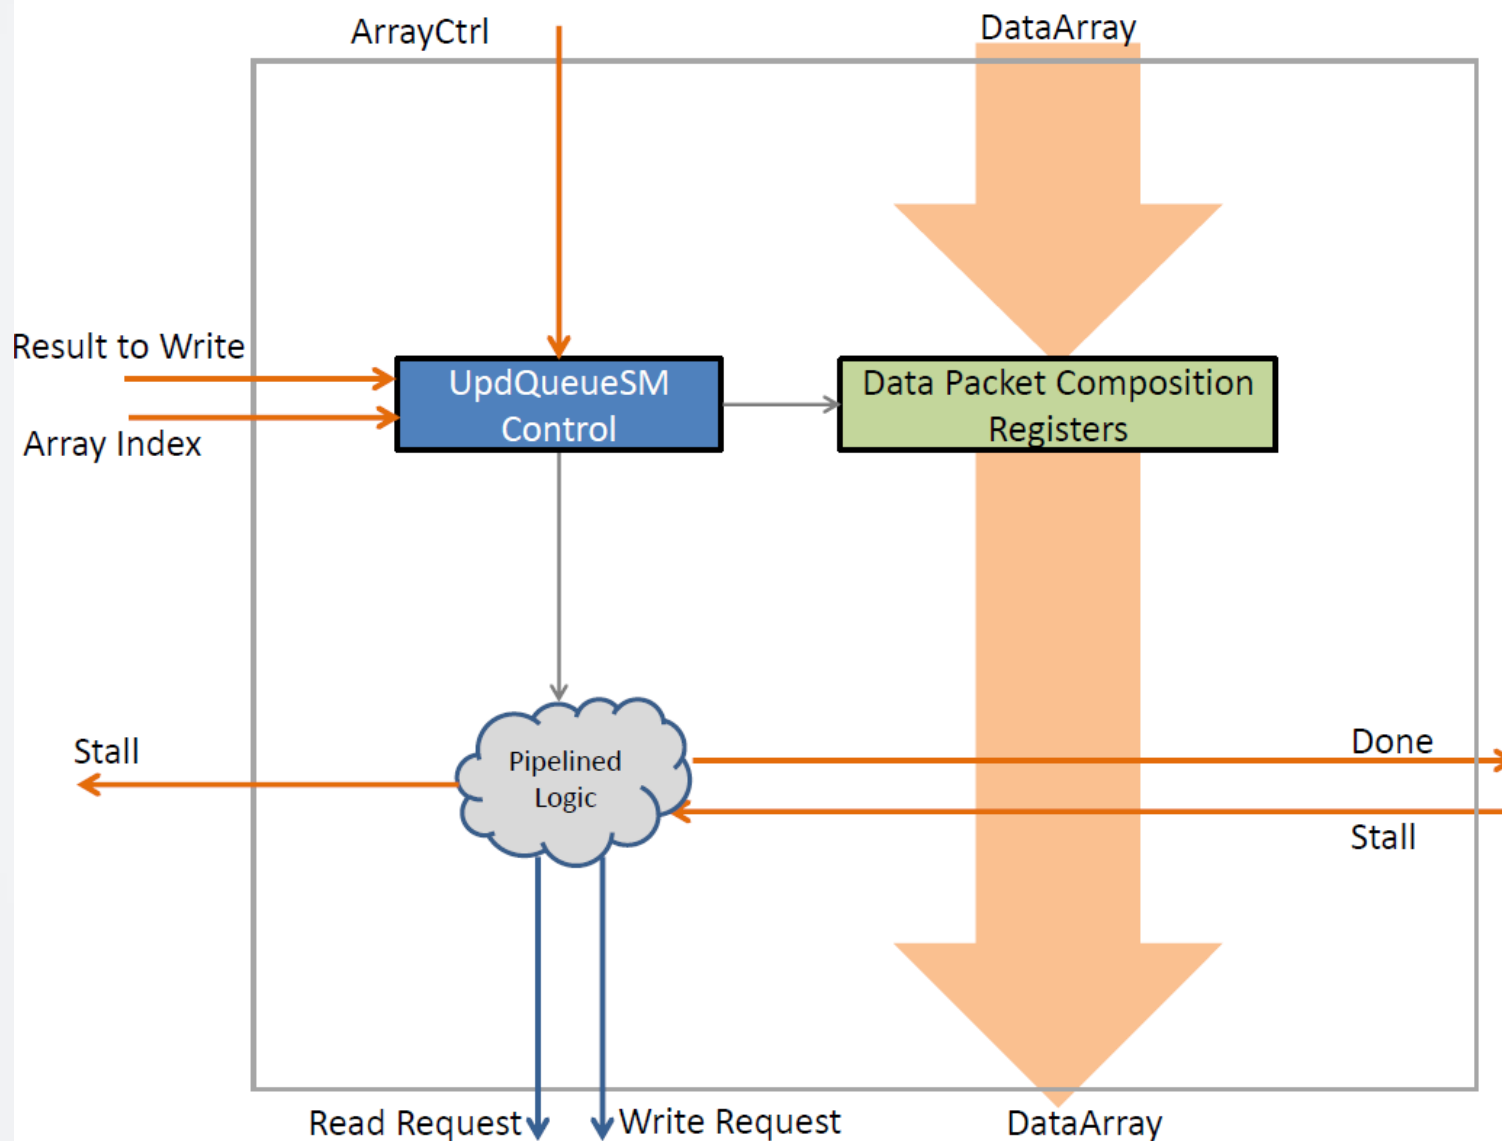
# The GraphOps Library: Utility Blocks



# Neighbor Property Reducer



# Element Update



# X-Stream: Streaming Graphs on CPUs

- Graph processing system using commodity hardware
- Sequentially streams entire edge lists, generates updates on active edges
- Designed to take advantage of sequential memory – absolutely no memory lookups necessary



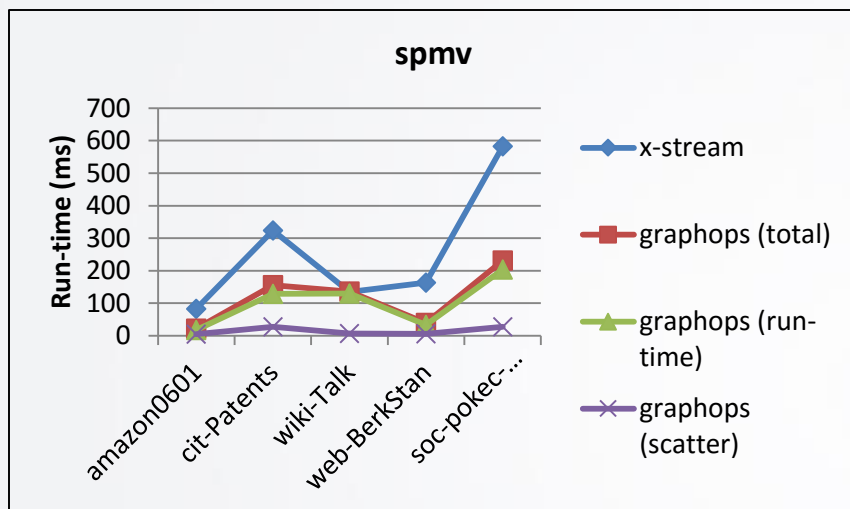
# X-Stream Comparison: Datasets

Datasets			
Name	Nodes	Edges	Description
amazon0601	475K	3.4M	Amazon product co-purchasing network from June 1 2003
cit-Patents	3.8M	16.5M	Citation network among US Patents
wiki-Talk	2.4M	5M	Wikipedia talk (communication) network
web-BerkStan	685K	7.6M	Web graph of Berkeley and Stanford
soc-Pokec	1.6M	30.6M	Pokec online social network

*Datasets are courtesy of the Stanford SNAP project. [snap.stanford.edu](http://snap.stanford.edu)*



# X-Stream Comparison



## Power Comparison

X-Stream: 190 W (2 sockets, TDP)

GraphOps: ~25 W

# Potential Future Work

- Higher level synthesis tool to target the GraphOps library
- Hide data transfer latency with double buffering and asynchronous execution
- Investigate locality-optimized storage for other sparse domains, e.g. machine learning
- Batch updates for host-side application
- Multi-FPGA
- Dynamic Graphs

# Memory Consistency

- Single writer per array
- If a GraphOps block is modifying an array, only that block may be simultaneously reading from the array
- Replicated arrays are read-only. Updates are made to the standard property array.
  - Use a SCATTER operation at the end of the computation

# Locality-Optimized Format: on CPUs

## ➤ PageRank

- 2M nodes, 16M edges

- OMP-C++, 4-thread

  - Current run-time: 3040

  - With replicated arrays: 1610

- Advantage was erased with the scatter

# Locality-Optimized Format: on CPUs

- Colleague (Chris) has been working on graph storage formats
- He attempted to implement my idea as part of a CPU run-time
- The scatter nullifies the advantage of the coalesced memory accesses

**Bandwidth Usage**

Category	Percentage
Rep Array (Read)	93%
Prop Array (Read)	2%
Prop Array (Write)	2%
Node Array (Read)	2%

Page Rank

**Methodology**

- Instrumented memory interface units with counters

**Line Bandwidth: 6.4 GB/s**

**Constraint Factors**

The diagram illustrates the Vertex-Element architecture. It consists of three main components: a Vertex Reader, a Nbr Reducer, and an Elem Updater, all connected to a shared DRAM (Dynamic Random Access Memory) at the bottom. The Vertex Reader reads data from the DRAM and sends it to the Nbr Reducer. The Nbr Reducer then sends data to the Elem Updater. The Elem Updater updates the data in the DRAM. The DRAM is labeled 'DRAM' and is represented by a green bar at the bottom. The components are labeled 'Vertex Reader', 'Nbr Reducer', and 'Elem Updater'.

The diagram illustrates four arrays used in a graph processing pipeline:

- Node Array:** A sequence of 8 nodes: 0, 1, 3, 6, 10, 13, 16, 18.
- Edge Array (not used):** A sequence of 16 edges: 2, 3, 4, 0, 5, 6, 1, 4, 5, 6, 1, 3, 5, 2, 3, 4, 2, 3.
- Property Array (Pagerank):** A sequence of 8 properties: p0, p1, p2, p3, p4, p5, p6, p7.
- duplicated Array (Pagerank):** A sequence of 16 properties: p2, p3, p4, p0, p5, p6, p1, p4, p5, p6, p1, p3, p5, p2, p3, p4, p2, p3. The first 10 elements are highlighted in red.

- Intel Xeon 5650 @ 2.7GHz
  - 2 sockets, 12 cores, 24 threads
  - **Bandwidth: 32 GB/s per socket**
  - **3 Memory Channels**
- FPGA: Xilinx Virtex-6 (150MHz)
  - Connected to host via PCIe8 Gen 2
  - **Bandwidth: 38.4 GB/s**

A stacked bar chart titled 'Page Rank' showing the distribution of page ranks for four categories: Prop Array (Write), Prop Array (Read), Rep Array (Read), and Node Array (Read). The y-axis represents the percentage from 0% to 120% in 20% increments. The x-axis is labeled 'Page Rank'. The bar is composed of four segments: a small blue segment at the bottom (2%), a large red segment (93%), a small green segment (2%), and a small purple segment at the top (2%).

Category	Page Rank (%)
Prop Array (Write)	2%
Prop Array (Read)	2%
Rep Array (Read)	93%
Node Array (Read)	2%

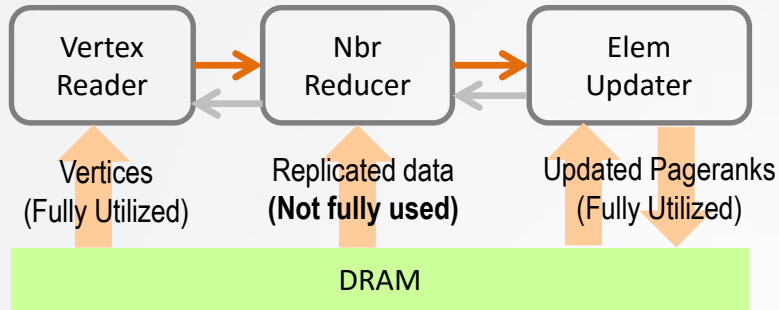
- Instrumented memory interface units with counters

### Constraining Factors:

- Locality
  - Optimal: Sequential access
  - Using: Alternating reads to the different arrays. All units operating simultaneously
- Packet size:
  - Optimal: 384 bytes x 4
  - Using: 192 bytes x 2
- Burst size:
  - Optimal: as large as possible (max 256)
  - Using: usually 1-2 (enough for nbr set)

# Bandwidth Efficiency

## Streaming Architecture: Page Rank Accelerator



### Usage Calculations

- Replicated array requests (number of bursts):  
1, 1, 1, 2, 1, 1, 1, 2, ...
- Average number of bursts per request: 1.22
  - Divided instrumented value of repl data divided by number of nodes
- Average number of bursts used per request: Assuming uniform with average degree of 8: 8 nbrs is 0.25 bursts. So usage rate is:  $0.25/1.22 = 0.205$
- Expected nbr bandwidth is:  $6.4 \text{ GB/s} * 0.205 = 1.312 \text{ GB/s}$
- Peak performance of PageRank is:  $36 \text{ MEPS} == 216 \text{ MB/s}$ 
  - About a factor of 1/6 of the expected performance
- Cause of performance being dropped on the floor: Single Memory Channel
  - Queuing/Switching: Nbr Reducer has to wait on the other requests using the memory channel concurrently and pay the cost of switching the active bank/rank/columns etc
- Ideally: multiple memory channels. One of them dedicated to Replicated data for streaming.

# Limitations of the GraphOps Library


- Limited expressability
- Limited portability
- Requires coalesced data for efficiency <sup>1</sup>
  - Common graph formats lead to highly inefficient memory behavior

<sup>1</sup> *Efficient Parallel Graph Exploration on Multi-core CPU and GPU.* Hong, Oguntebi, et al.



# Streaming Processors

- Multiple “functional units” execute simultaneously
- Each function performs a different task on the data stream flowing through it
- GraphOps blocks are implemented as coarse-grained functions
- More simple approach for end user: higher level building blocks



# Disadvantages to Graph Replication

# Additional GraphOps diagrams

# Figures from FPGA Paper

# GraphOps are Parameterizable

# Edge Properties

- A logical way of describing the locality-optimized format
  - Think of the data as being associated with an edge instead of a vertex

# Approach

- Initially started with a domain-specific HLS approach
- Was hoping to build full applications on the FPGA
  - Sensitive control was difficult / time-consuming to generate automatically in hardware
  - Especially without an ISA and full architecture
  - Memory behavior was bad anyway
- Converted to an accelerator-based approach

# Real-world Dataset properties

- (from snap website)



# How Different from GPUs and CPU Vector Machines

# Scatter/Gather Options in HW

# About Us

We develop brand name with individual creative solutions and help our customers to **ear money**

## BUSINESS ANALYTICS

Architecto beatae vitae dicta  
sunt explicabo nemo enim  
ipsam voluptatem. Architecto  
beatae vitae dicta.

**Lorem Ipsum** is simply dummy text of the printing and typesetting industry. Lorem Ipsum has been the industry's standard dummy text ever since the 1500s, when an unknown printer took a galley of type and scrambled it to make a type specimen book. It has survived not only five centuries, but also the leap into electronic typesetting, remaining essentially unchanged.

# Team Work Sample

We develop brand name with individual creative solutions and help our customers to **ear money**




**Marcus Lopez**  
*Designer*

Architecto beatae vitae  
dicta sunt explicabo nemo  
enim ipsam voluptatem.




**María Castro**  
*Marketing*

Architecto beatae vitae  
dicta sunt explicabo nemo  
enim ipsam voluptatem.



**Carlos Perez**  
*Animation*

Architecto beatae vitae  
dicta sunt explicabo nemo  
enim ipsam voluptatem.

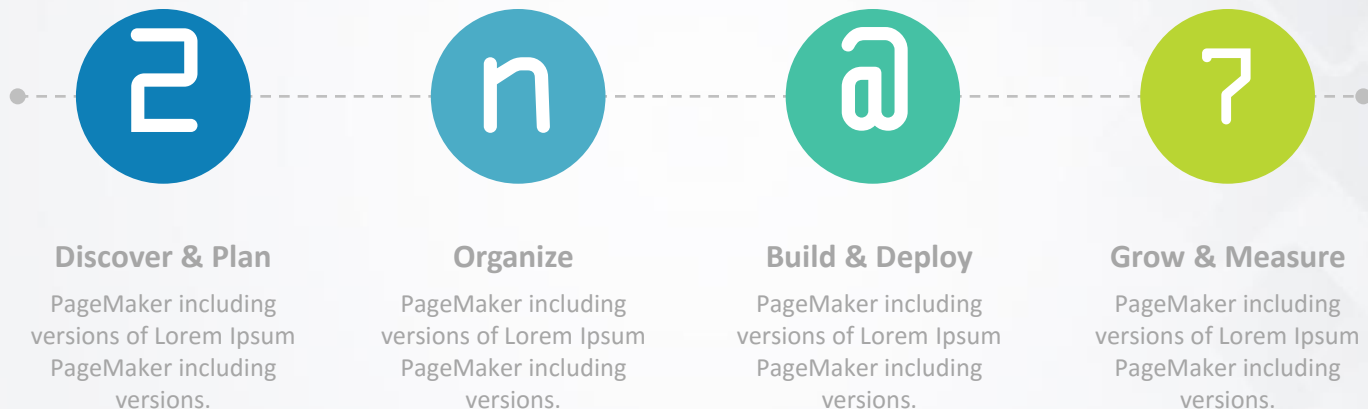


**Antonio Ruiz**  
*Sales Rep*

Architecto beatae vitae  
dicta sunt explicabo nemo  
enim ipsam voluptatem.

# Work Process Sample

Contrary to popular. It has roots in a piece of classical Latin **our process, from start to finish.**



**Lorem Ipsum** is simply dummy text of the printing and typesetting industry. Lorem Ipsum has been the industry's standard dummy text ever since the 1500s, when an unknown.

# Services List Sample

We develop brand name with individual creative solutions and help our customers to **ear money**



## INVESTMENT CONSULTING

Architecto beatae vitae dicta  
sunt explicabo nemo enim  
ipsam voluptatem. Architecto  
beatae vitae dicta.

## TAX STRATEGIES

Architecto beatae vitae dicta  
sunt explicabo nemo enim  
ipsam voluptatem. Architecto  
beatae vitae dicta.

## BROKER COMPARISION

Architecto beatae vitae dicta  
sunt explicabo nemo enim  
ipsam voluptatem. Architecto  
beatae vitae dicta.

## BUSINESS ANALYTICS

Architecto beatae vitae dicta  
sunt explicabo nemo enim  
ipsam voluptatem. Architecto  
beatae vitae dicta.

# 3 Columns Sample

We develop brand name with individual creative solutions and help our customers to **ear money**

**Lorem Ipsum** is simply dummy text of the printing and typesetting industry. Lorem Ipsum has been the industry's standard dummy text ever since the 1500s,

when an unknown printer took a galley of type and scrambled it to make a type specimen book. It has survived not only five centuries, but also the leap into electronic typesetting, remaining essentially unchanged.

**Lorem Ipsum** is simply dummy text of the printing and typesetting industry. Lorem Ipsum has been the industry's standard dummy text ever since the 1500s, when an unknown printer took a galley of type and scrambled it to make a type specimen book. It has survived not only five centuries, but also the leap into electronic typesetting, remaining essentially unchanged.

**Lorem Ipsum** is simply dummy text

of the printing and typesetting industry. Lorem Ipsum has been the industry's standard dummy text ever since the 1500s, when an unknown printer took a galley of type and scrambled it to make a type specimen book.

**Lorem Ipsum** is simply dummy text of the printing and typesetting industry. Lorem Ipsum has been the industry's standard dummy text ever since the 1500s, when an unknown printer took a galley of type and scrambled it to make a type specimen book.

## 2 Columns Sample

We develop brand name with individual creative solutions and help our customers to **ear money**

**Lorem Ipsum** is simply dummy text of the printing and typesetting industry. Lorem Ipsum has been the industry's standard dummy text ever since the 1500s, when an unknown printer took a galley of type and scrambled it to make a type specimen book. It has survived not only five centuries, but also the leap into electronic typesetting, remaining essentially unchanged.

**Lorem Ipsum** is simply dummy text of the printing and typesetting industry. Lorem Ipsum has been the industry's standard dummy text ever since the 1500s, when an unknown printer took a galley of type and scrambled it to make a type specimen book. It has survived not only five

**Lorem Ipsum** is simply dummy text of the printing and typesetting industry. Lorem Ipsum has been the industry's standard dummy text ever since the 1500s, when an unknown printer took a galley of type and scrambled it to make a type specimen book.

centuries, but also the leap into electronic typesetting, remaining essentially unchanged.

**Lorem Ipsum** is simply dummy text of the printing and typesetting industry. Lorem Ipsum has been the industry's standard dummy text ever since the 1500s, when an unknown printer took a galley of type and scrambled it to make a type specimen book. It has survived not only five centuries, but also the leap into electronic typesetting, remaining essentially unchanged.