

LMC: Automatic Resource-Aware Program-Optimized Memory Partitioning

**Hsin-Jung Yang[†], Kermin E. Fleming[‡], Michael Adler[‡],
Felix Winterstein[§], and Joel Emer[†]**

[†] Massachusetts Institute of Technology,

[‡] Intel Corporation, [§] Imperial College London,

February 22nd, FPGA 2016

Motivation



Moore's Law continues

- More transistors & memory controllers on modern FPGAs
 - Example: **Xilinx VC709**: two 4GB DDR3 memories
Nallatech 510T: eight 4GB DDR4 memories + 2GB HMC
Xeon + FPGA: three memory channels



It is difficult to fully utilize DRAM bandwidth

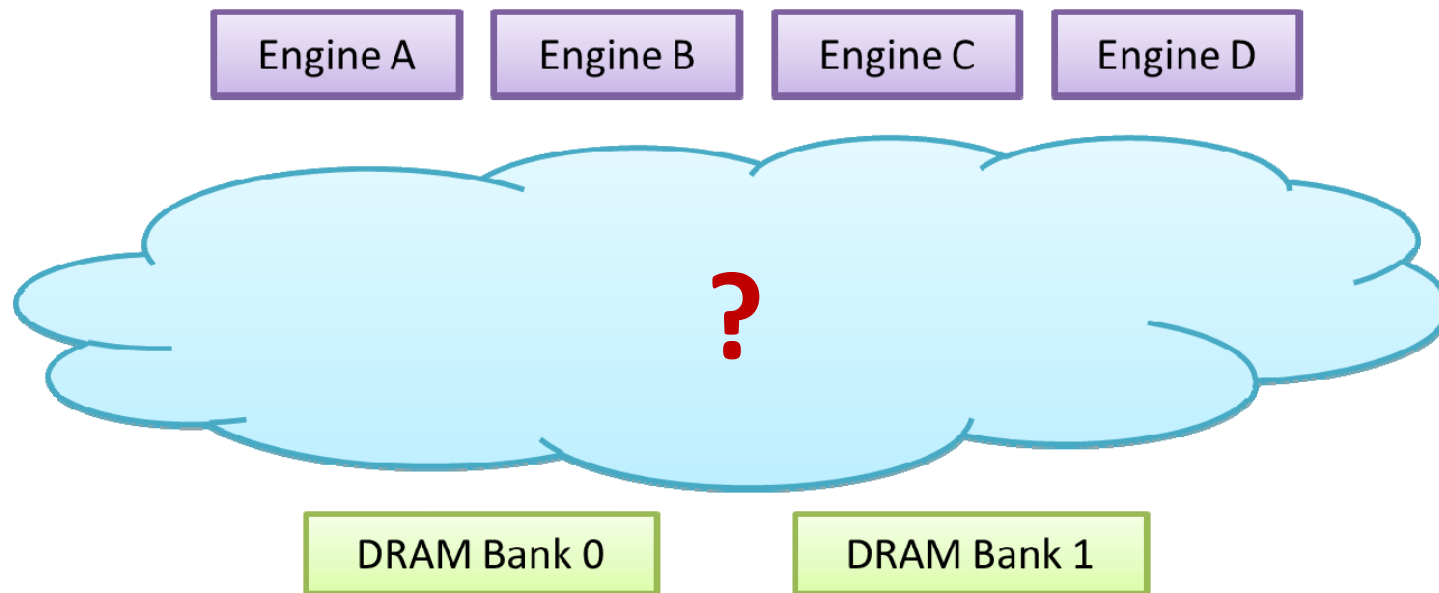
- Co-optimizing application cores and memory systems
- Porting an existing design to a new platform
 - Smaller FPGA -> Larger FPGA
 - Single FPGA -> Multiple FPGAs

Goal: Automatically optimizing the memory system to efficiently utilize the increased DRAM bandwidth

Utilizing Multiple DRAMs

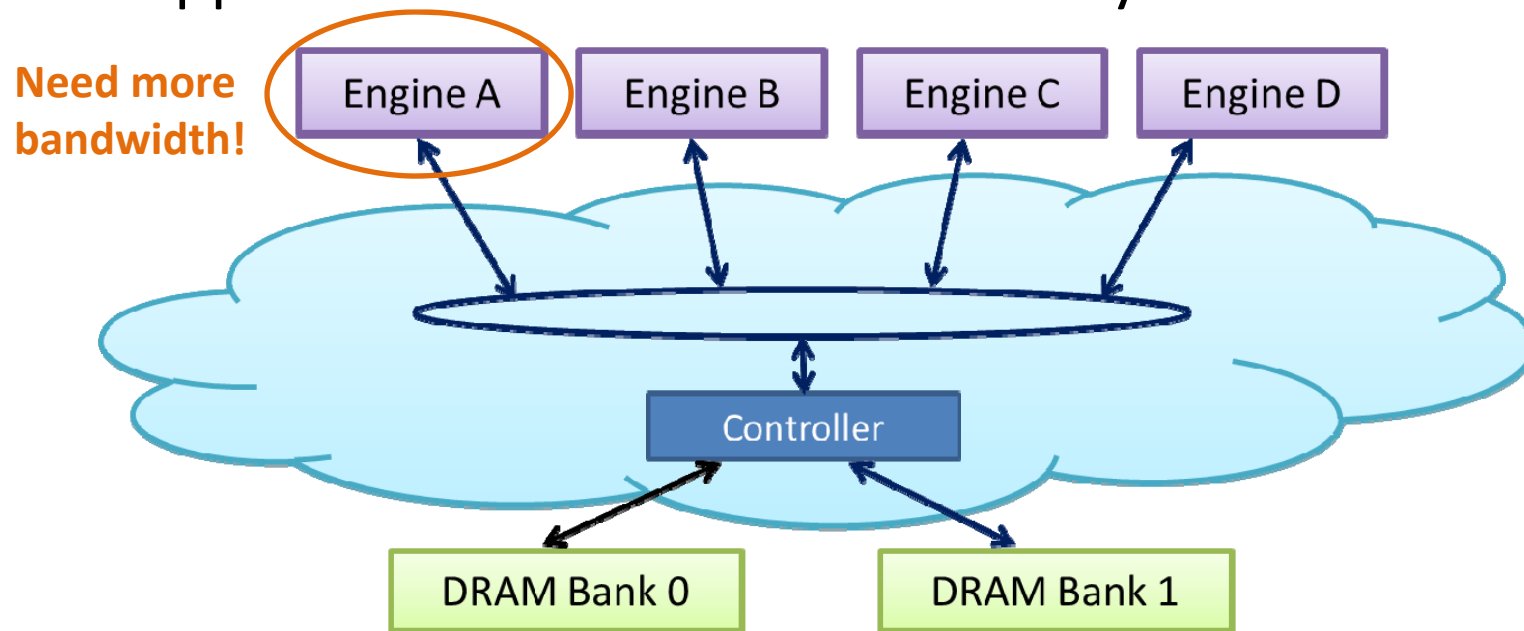
- How to connect computational engines to DRAMs in order to maximize program performance?
 - Network topology: latency, bandwidth
 - On-chip caching
 - Area constraints

High design complexity!



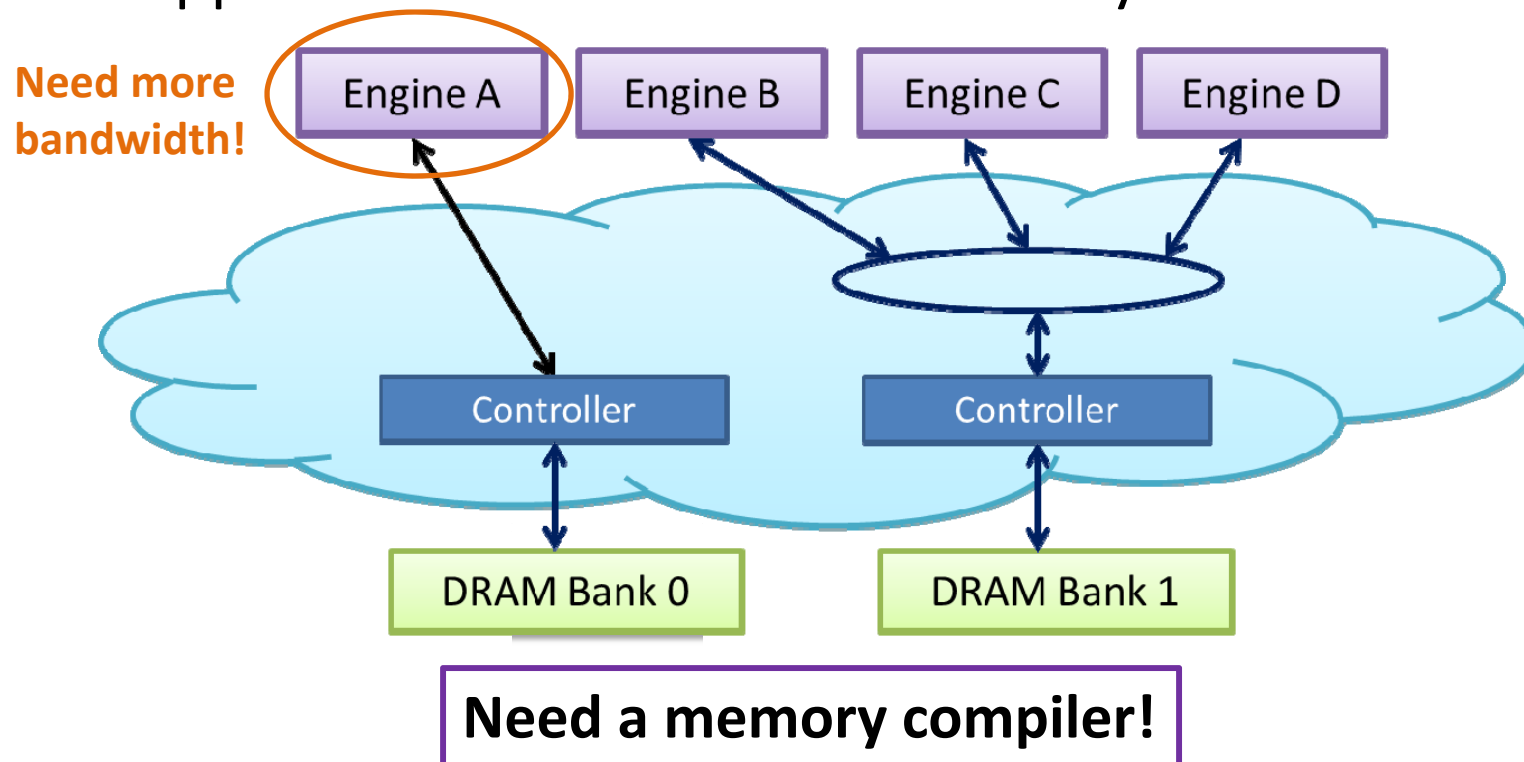
Utilizing Multiple DRAMs

- How to connect computational engines to DRAMs in order to maximize program performance?
 - High design complexity: network, caching...
- Applications have different memory behavior



Utilizing Multiple DRAMs

- How to connect computational engines to DRAMs in order to maximize program performance?
 - High design complexity: network, caching...
- Applications have different memory behavior

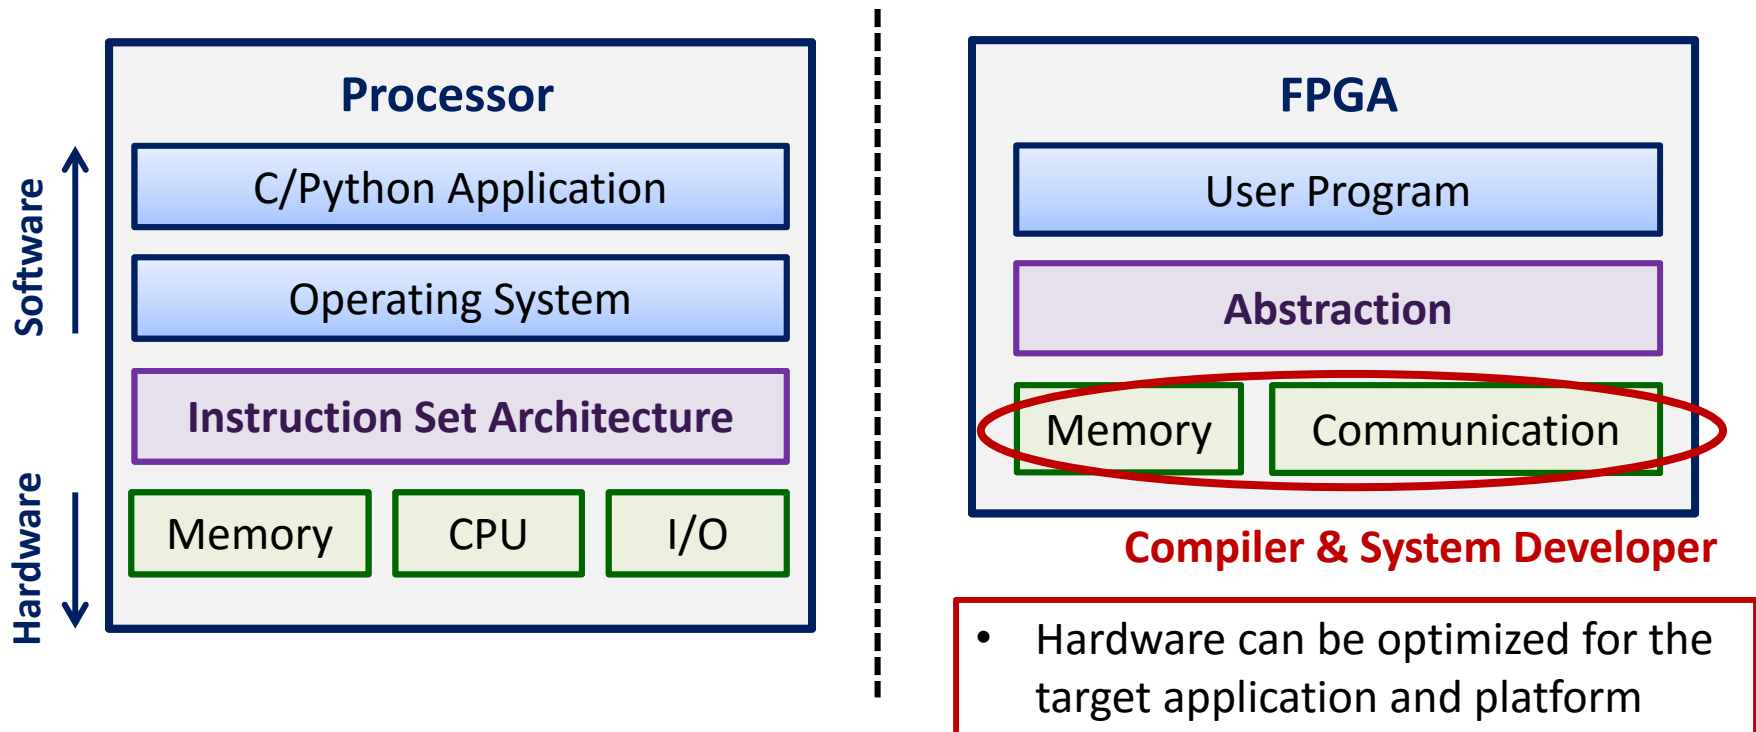


Automatic Construction of Program-Optimized Memories

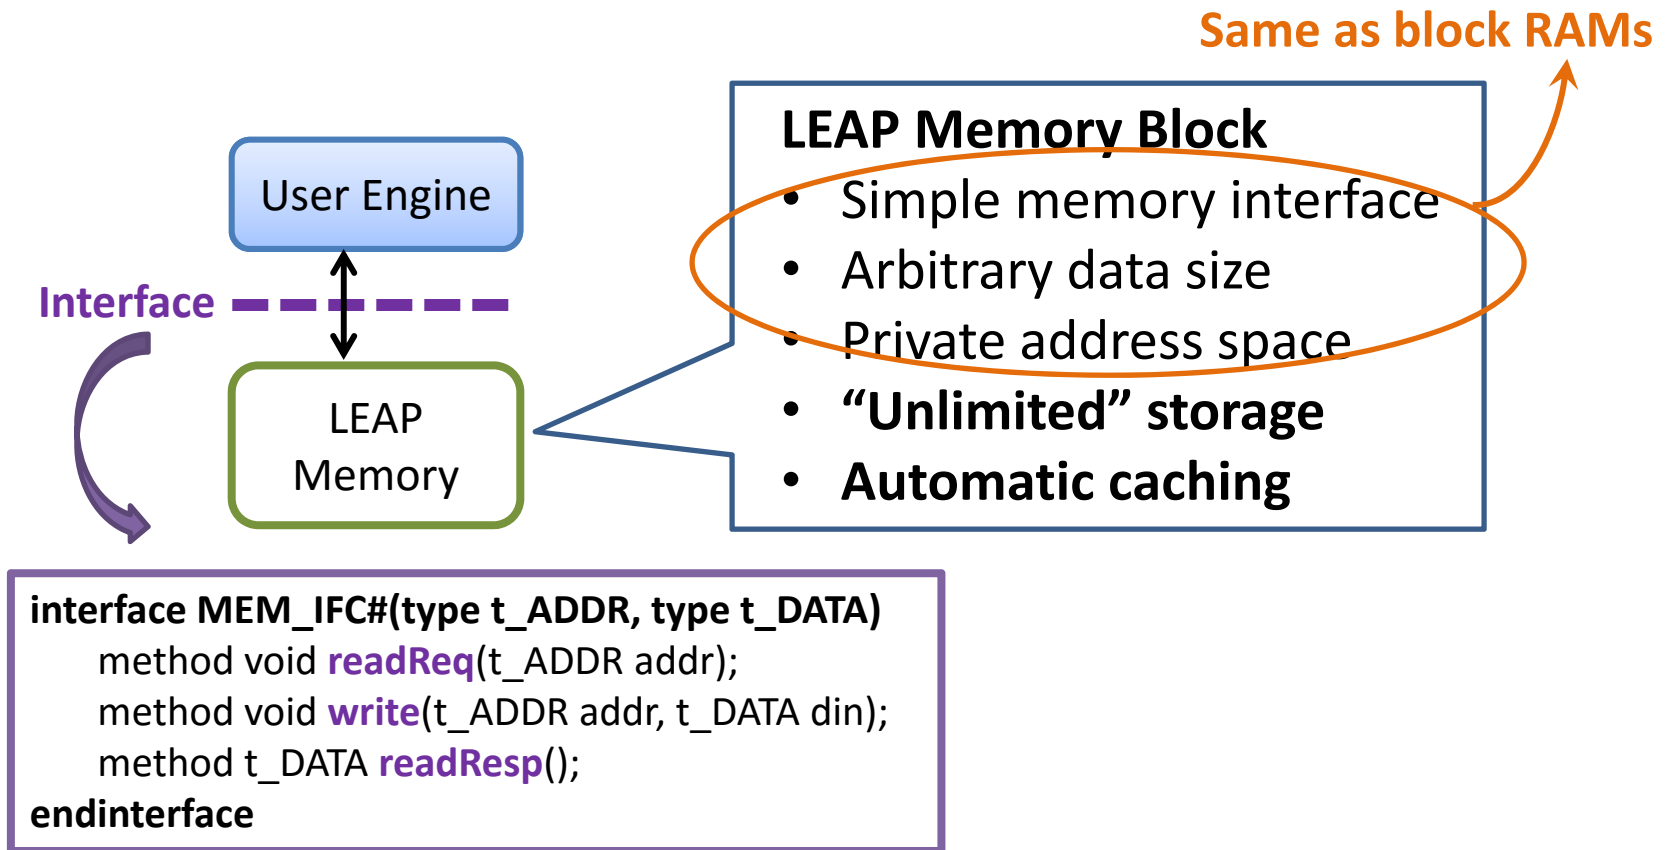
- **A clearly-defined, generic memory abstraction**
 - Separate the user program from the memory system implementation
- **Program introspection**
 - To understand programs' memory behavior
- **A resource-aware, feedback-driven memory compiler**
 - Use introspection results as feedback to automatically construct the “best” memory system for the target program and platform

Abstraction

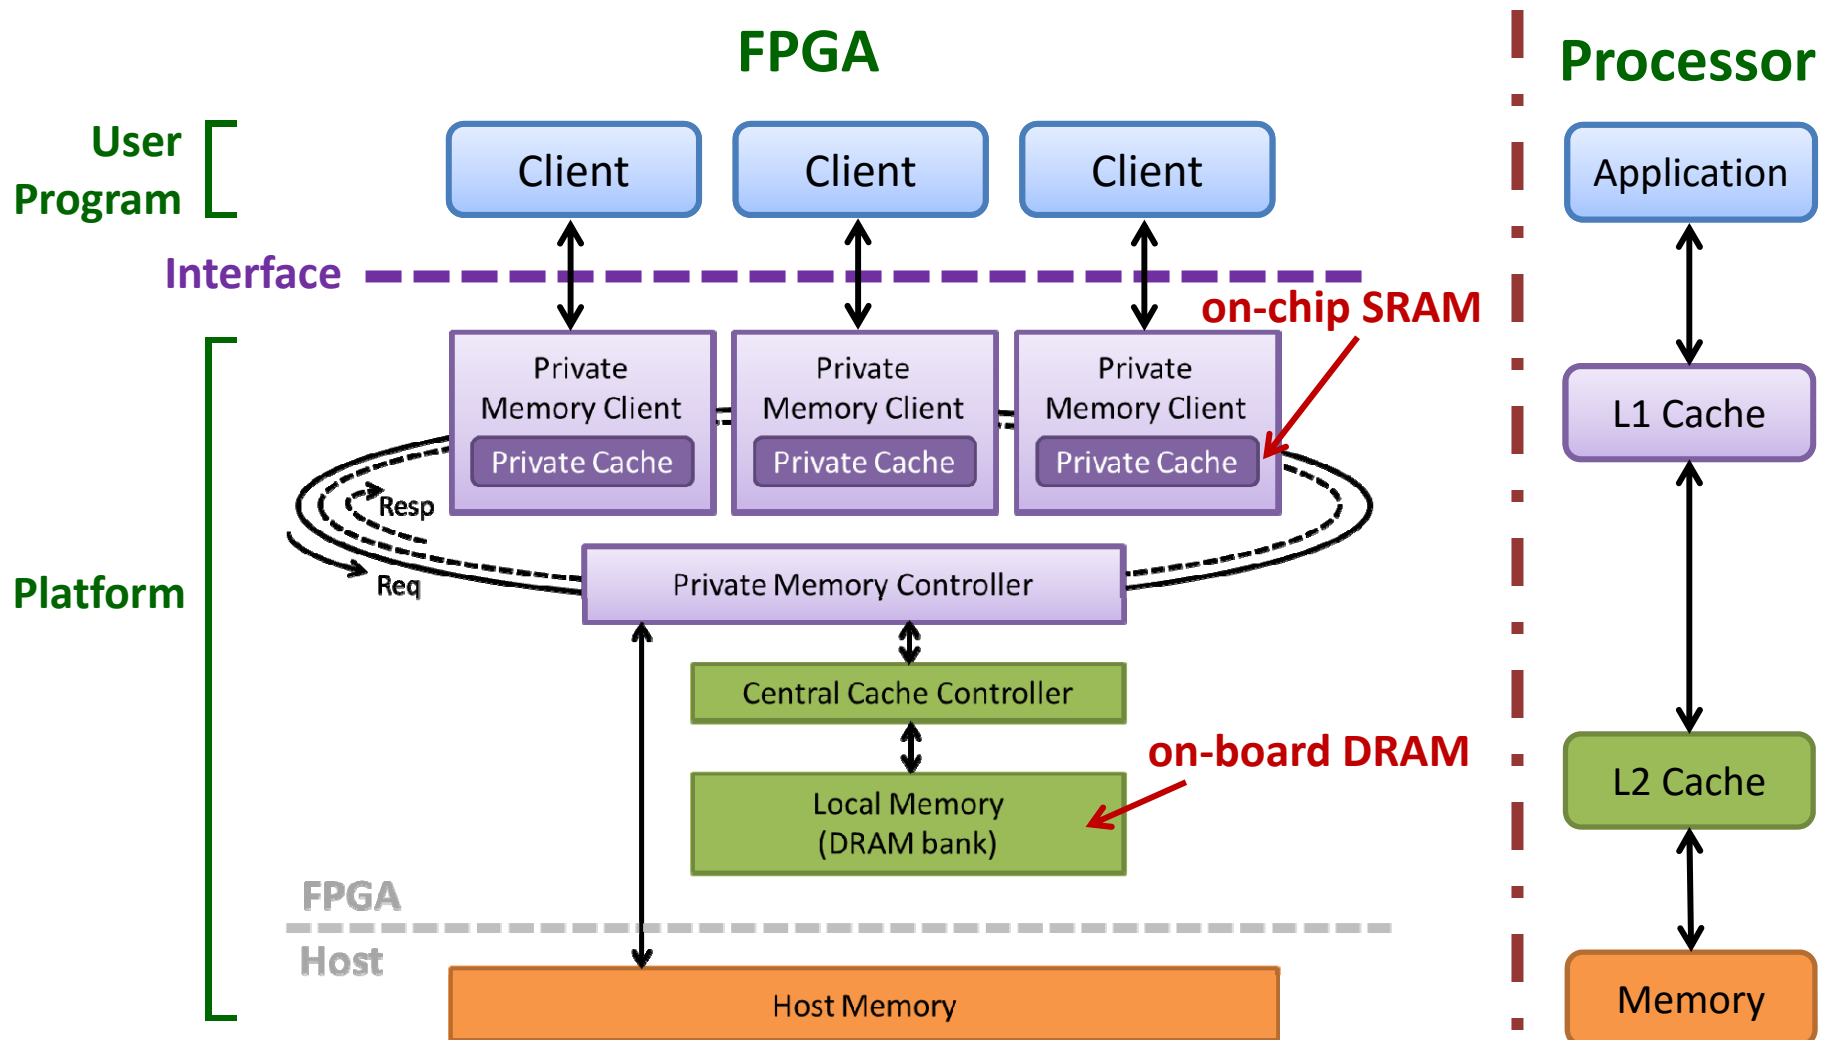
- Abstraction hides implementation details and provides good programmability



LEAP Memory Abstraction




LEAP Private Memory



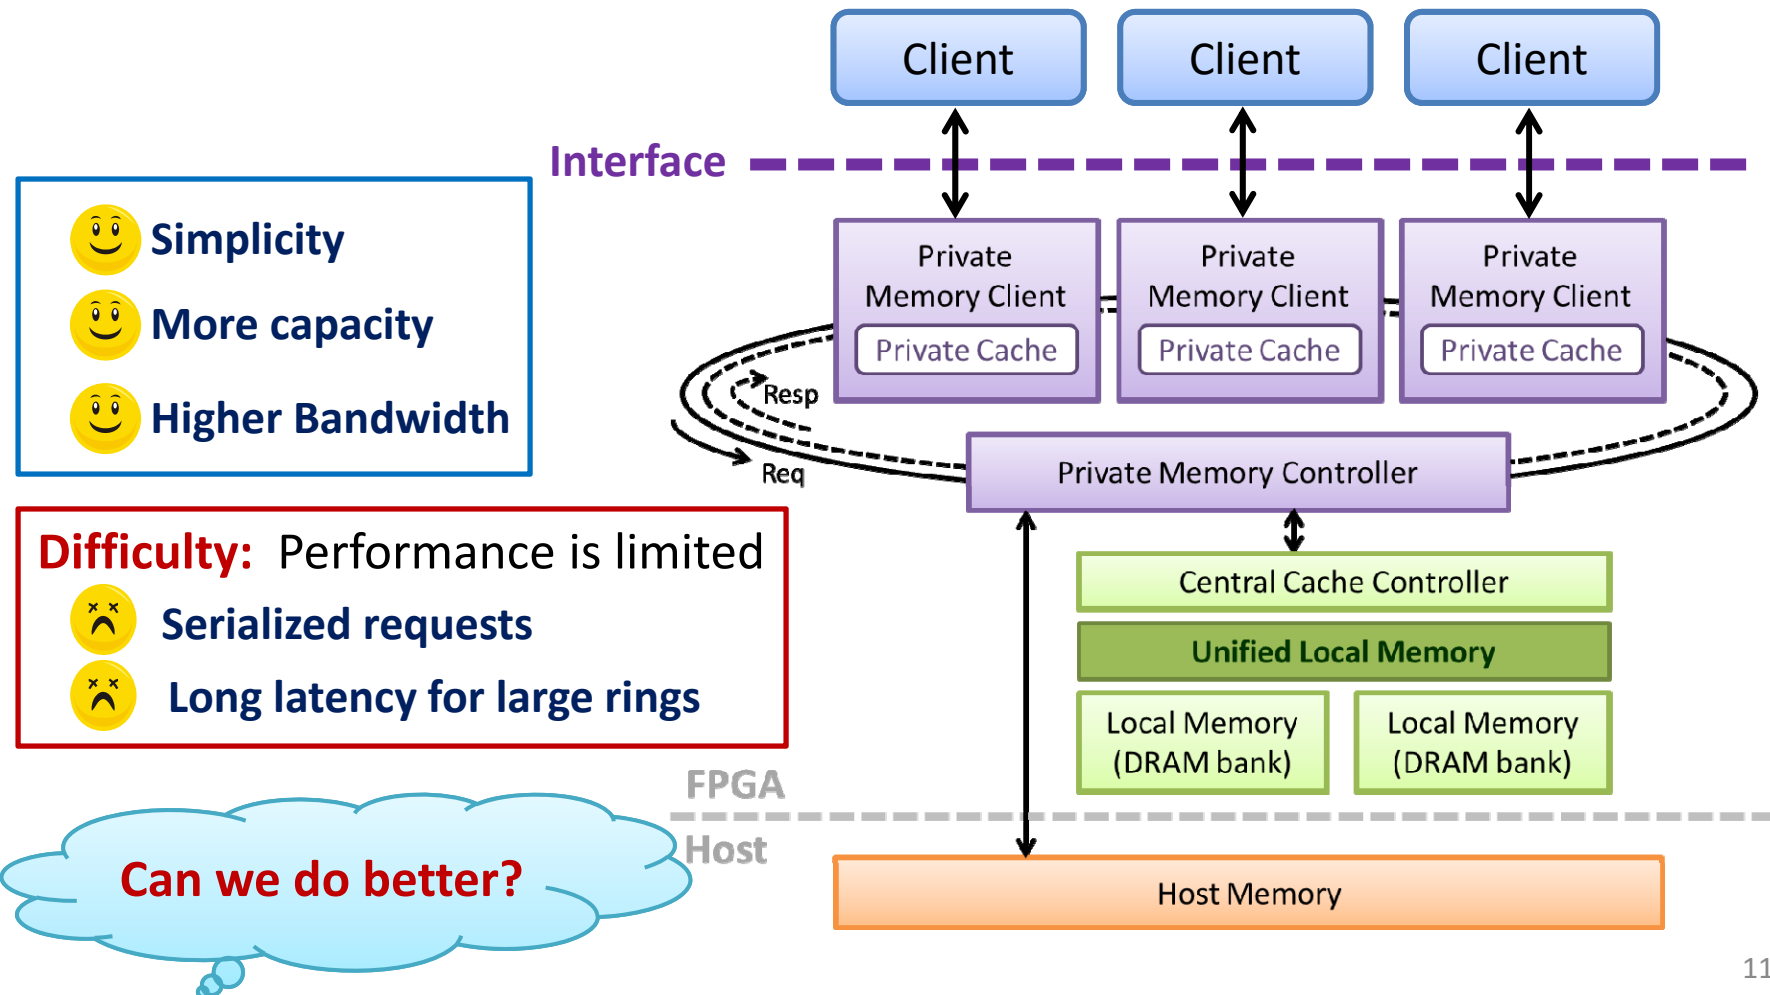
LEAP Memory with Multiple DRAMs

- **Naïve solution:** unified memory with multiple DRAM banks



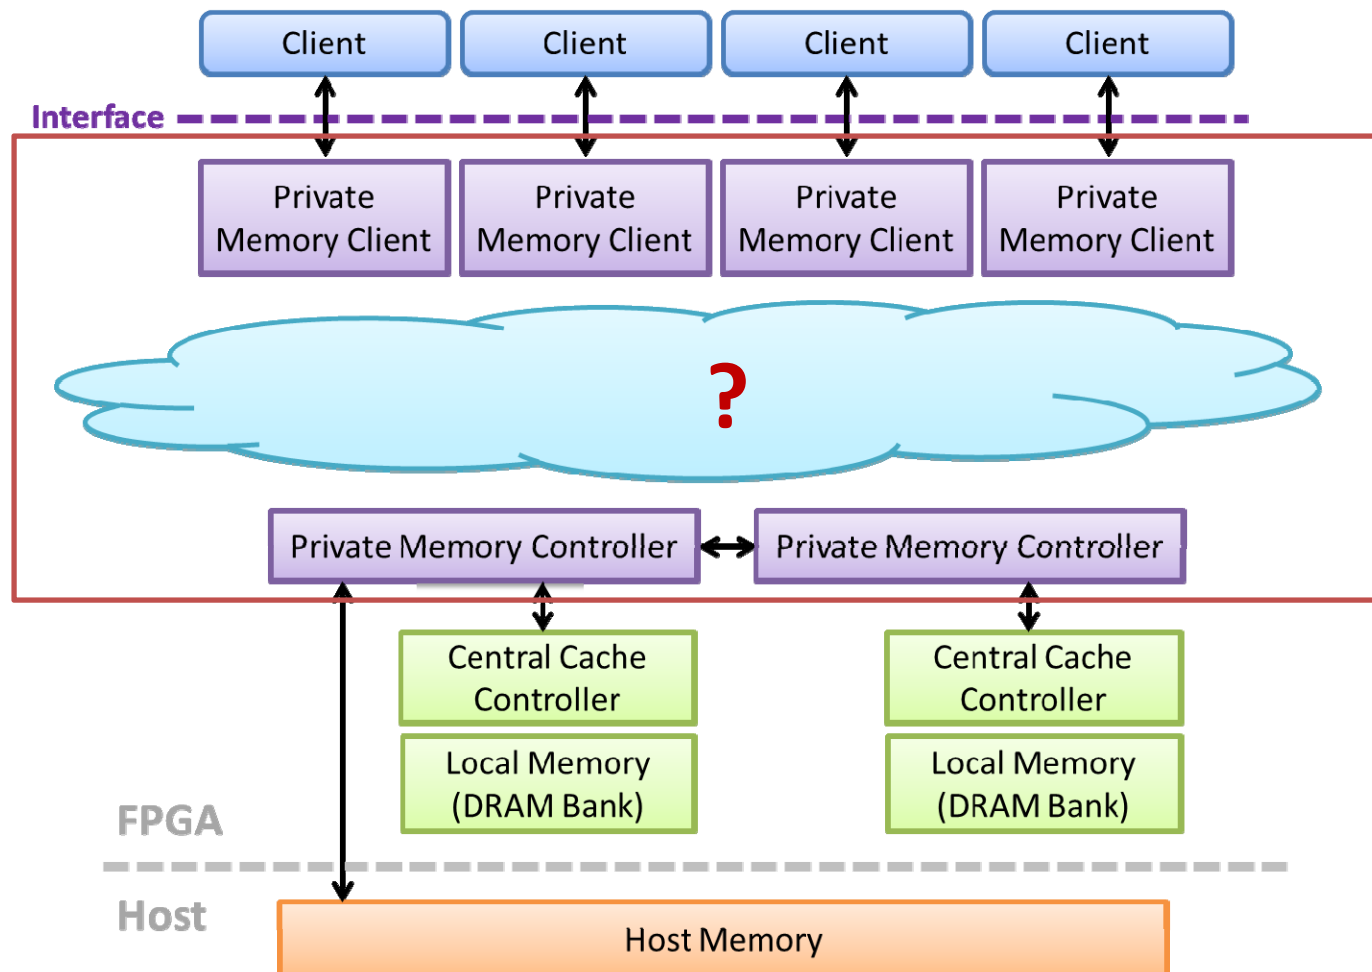
LEAP Memory with Multiple DRAMs

- **Naïve solution:** unified memory with multiple DRAM banks



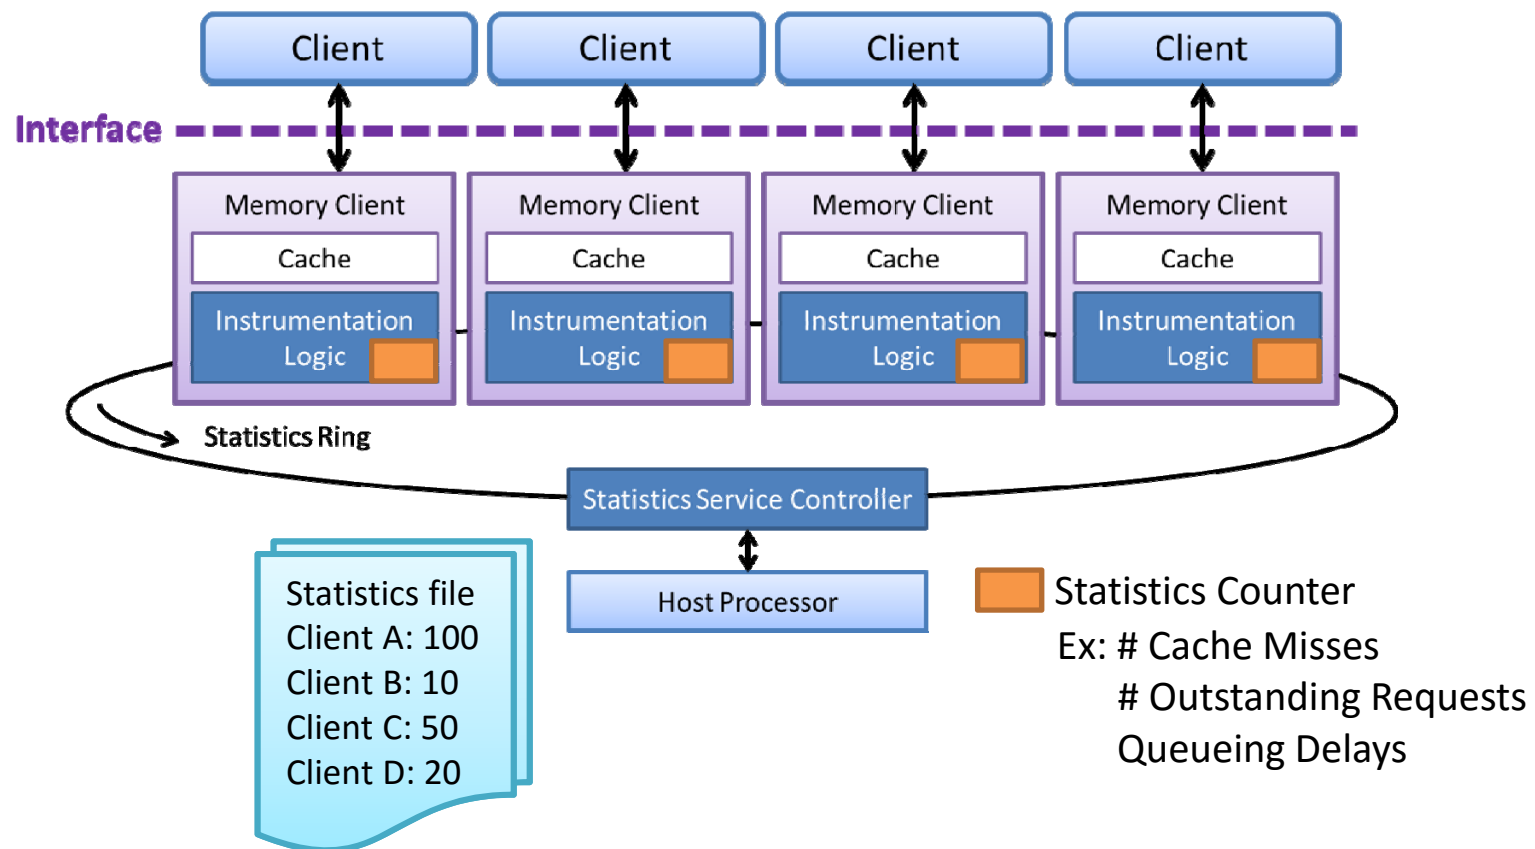
LEAP Memory with Multiple DRAMs

- **Distributed central caches and memory controllers**



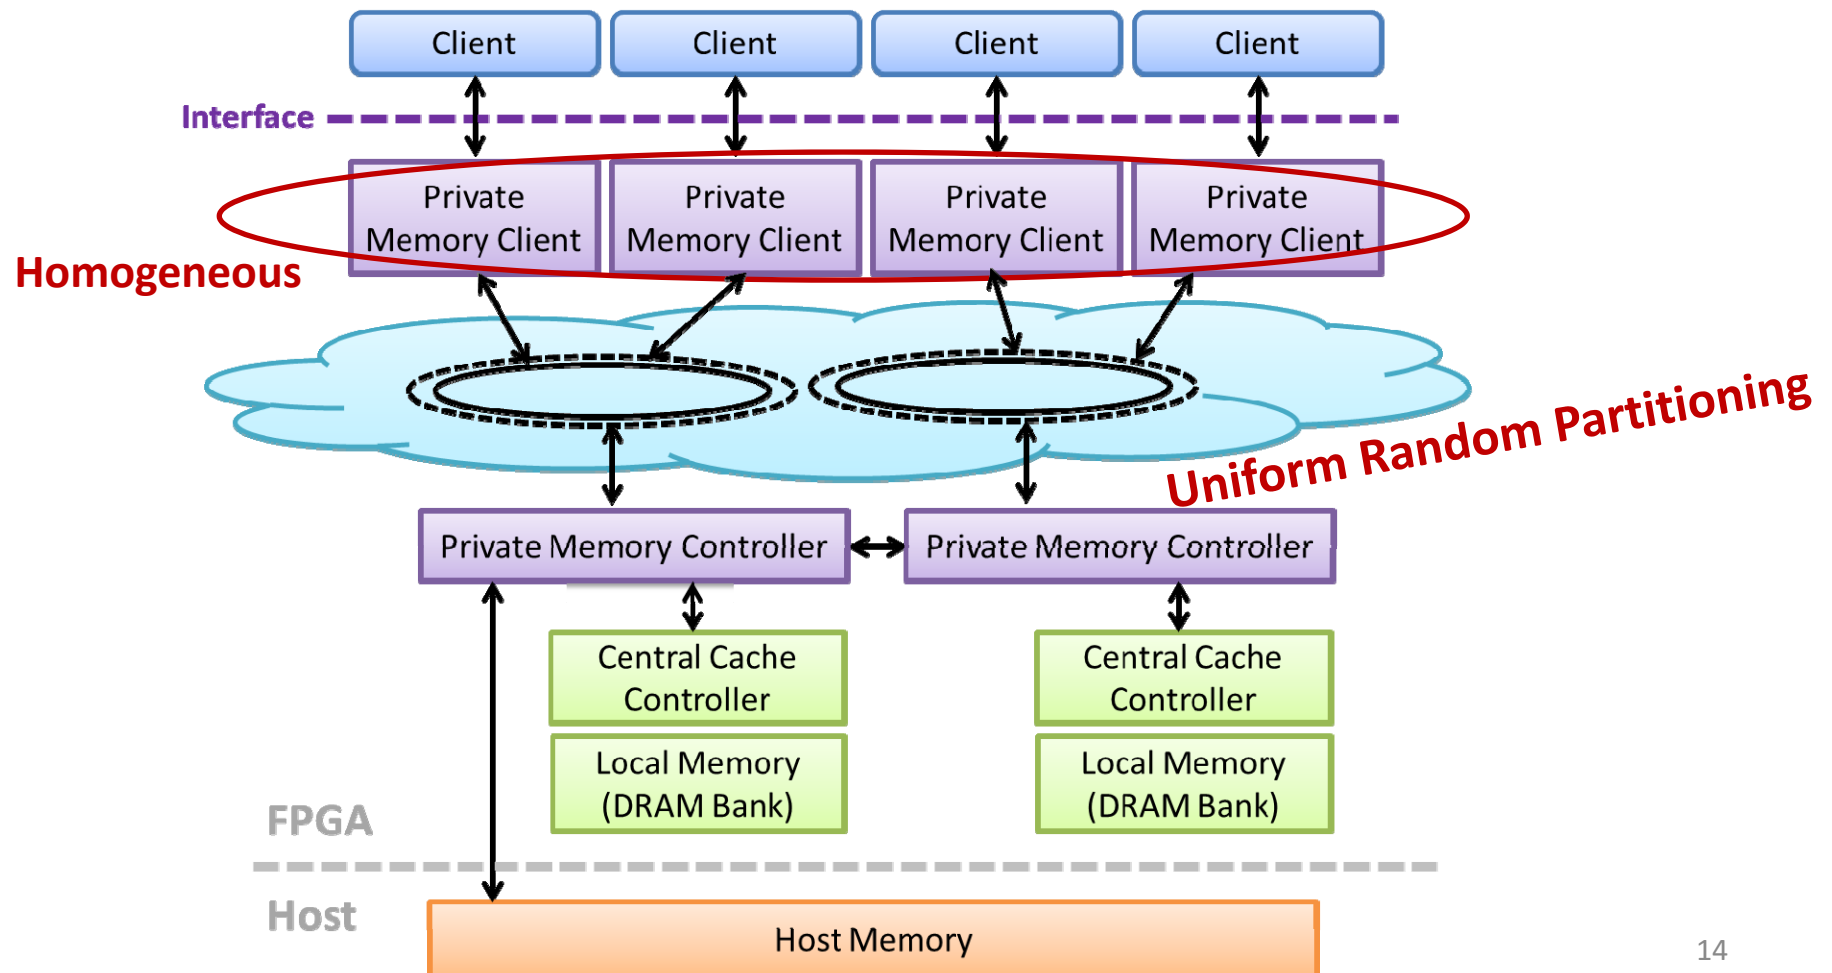
Private Cache Network Partitioning

- **Program introspection**
 - To understand programs' memory behavior



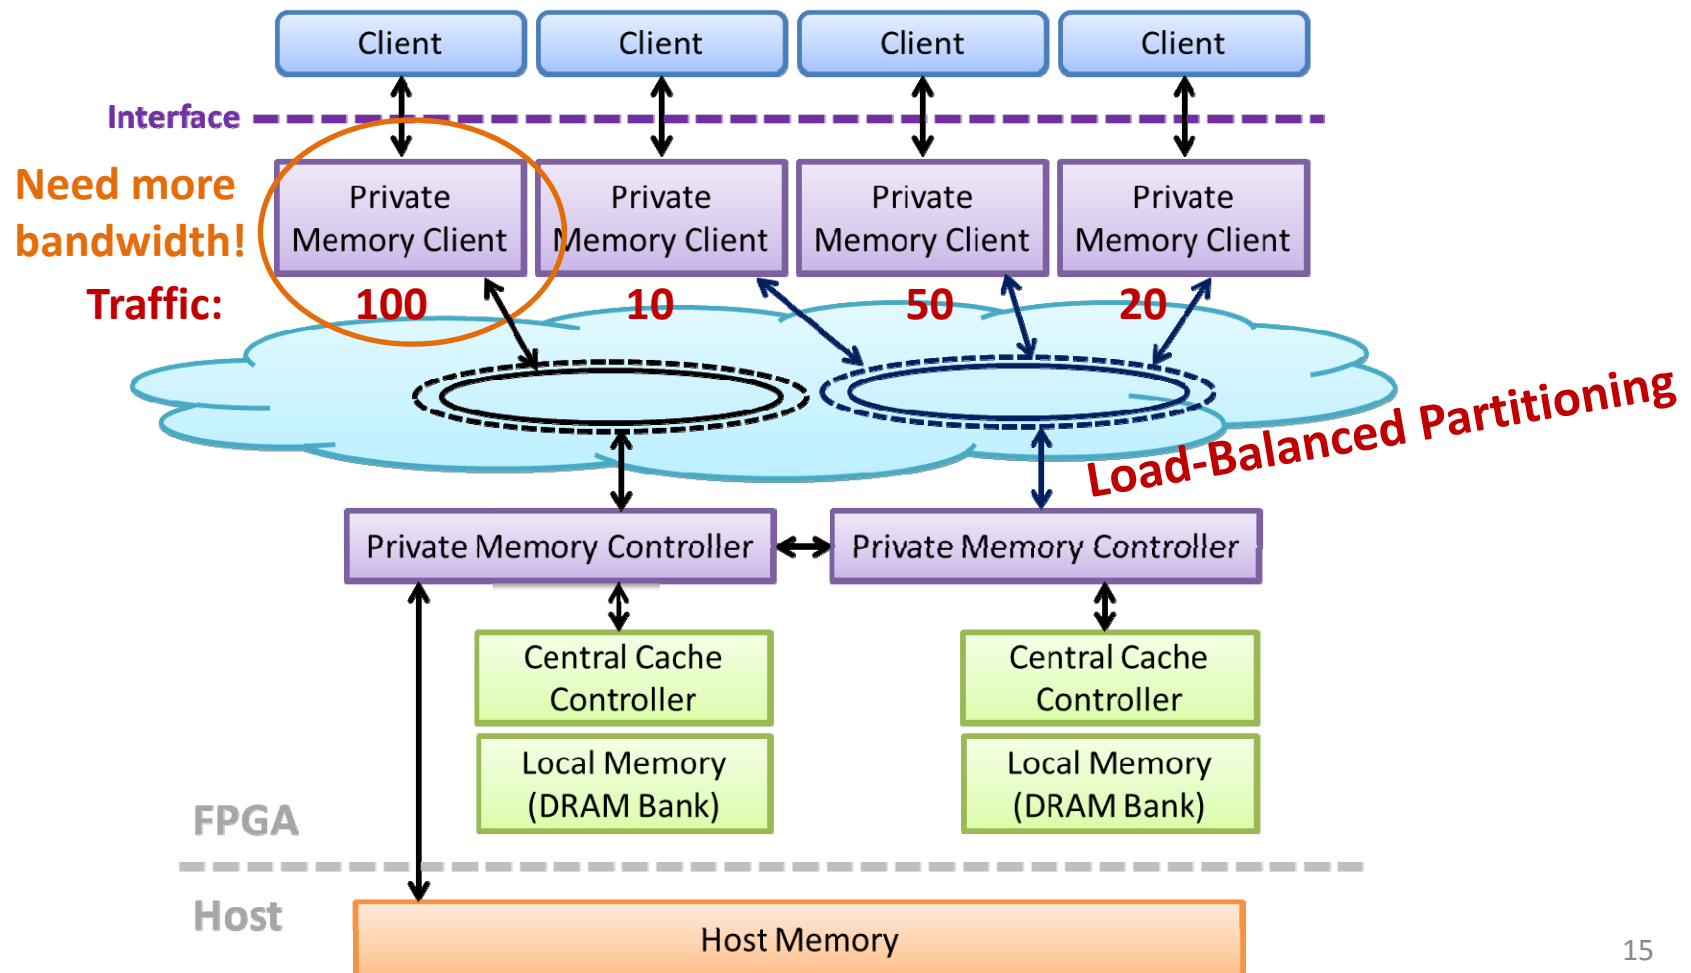
Private Cache Network Partitioning

- **Case 1: Memory clients with homogeneous behavior**



Private Cache Network Partitioning

- Case 2: Memory clients with heterogeneous behavior



Private Cache Network Partitioning

- **Case 2: Memory clients with heterogeneous behavior**
 - Load-balanced partitioning
 - Classical minimum makespan scheduling problem

m controllers, n clients, client j with traffic t_j

$$x_{i,j} = \begin{cases} 1 & \text{if client } j \text{ is mapped to controller } i \\ 0 & \text{otherwise} \end{cases}$$

ILP formulation:

minimize t

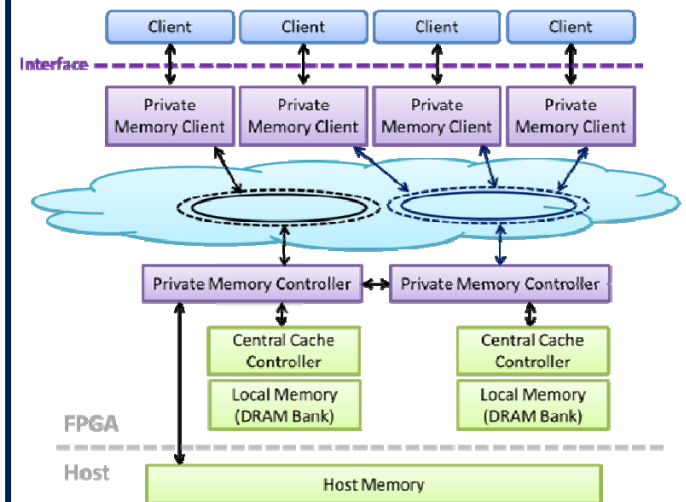
$$\text{s.t.} \quad \sum_{j=1}^n x_{i,j} t_j \leq t, \quad i = 1, \dots, m$$

$$\sum_{i=1}^m x_{i,j} = 1, \quad j = 1, \dots, n$$

$$x_{i,j} \in \{0,1\}, \quad i = 1, \dots, m, j = 1, \dots, n$$

Approximation:

Longest Processing Time (LPT)
Algorithm

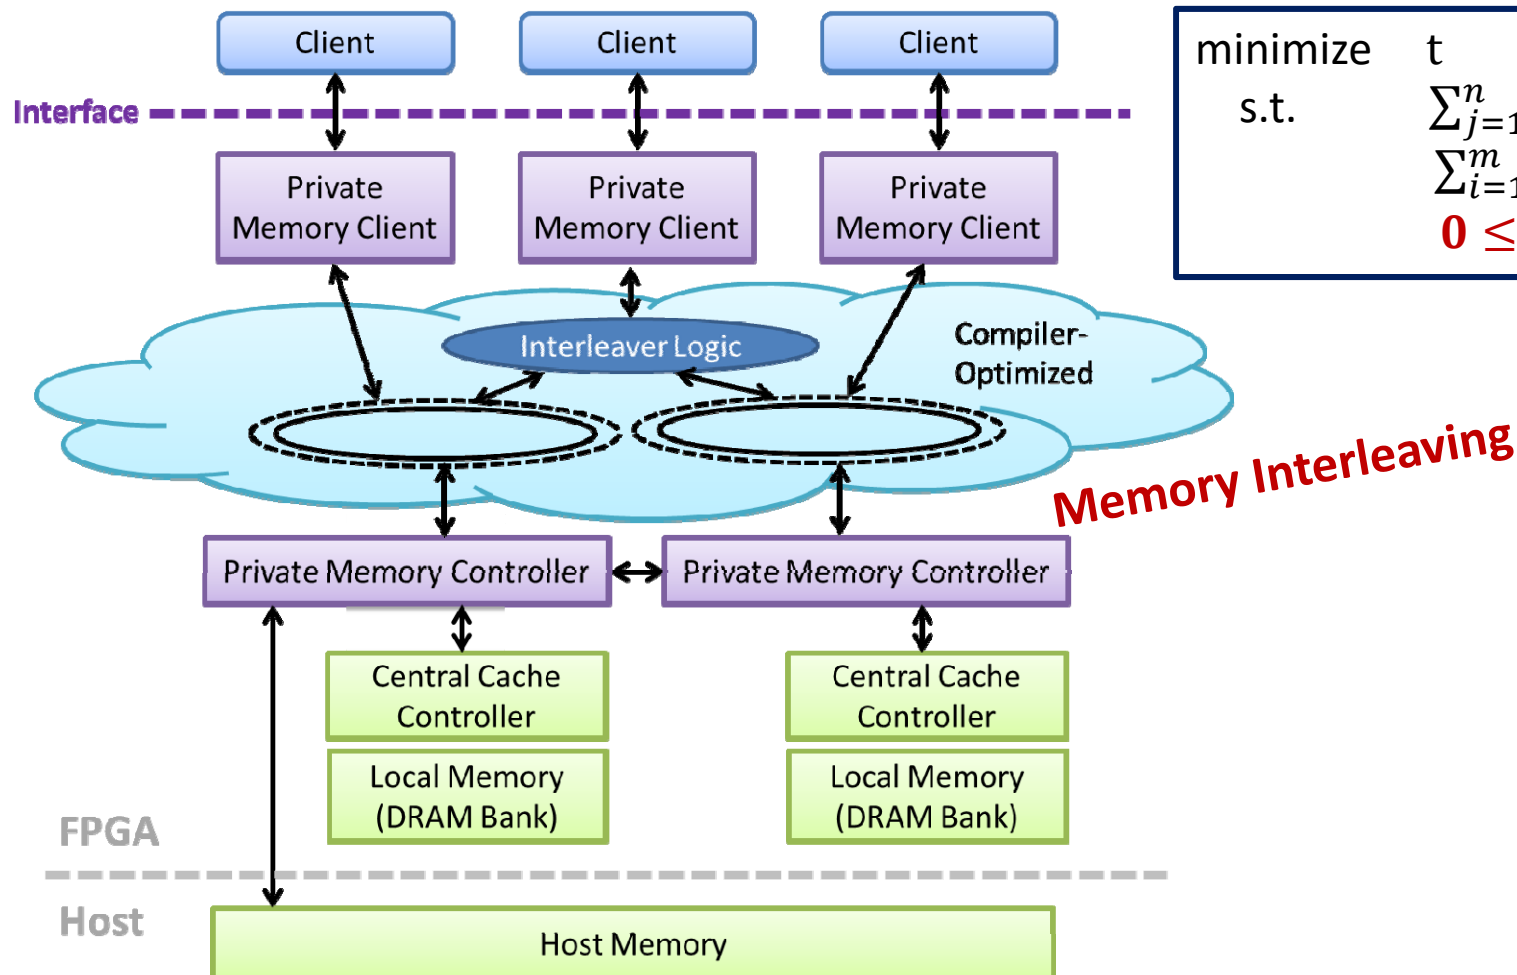


Private Cache Network Partitioning

- Case 3: Fractional load-balancing

ILP->LP

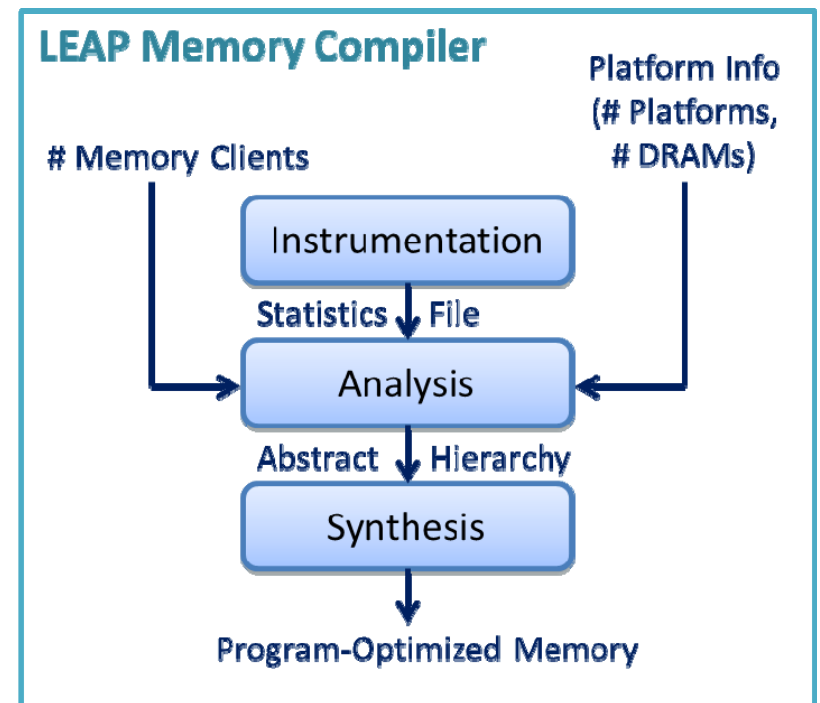
$$\begin{aligned} &\text{minimize} && t \\ &\text{s.t.} && \sum_{j=1}^n x_{i,j} t_j \leq t \\ & && \sum_{i=1}^m x_{i,j} = 1 \\ & && 0 \leq x_{i,j} \leq 1 \end{aligned}$$



LEAP Memory Compiler


- **Three-phase feedback-driven compilation**

- **Instrumentation (optional):** to collect runtime information about the way the program uses memory
- **Analysis:** to analyze the program properties and decide an optimized memory hierarchy
- **Synthesis:** to implement the program-optimized memory




LEAP Memory Performance

- Baseline



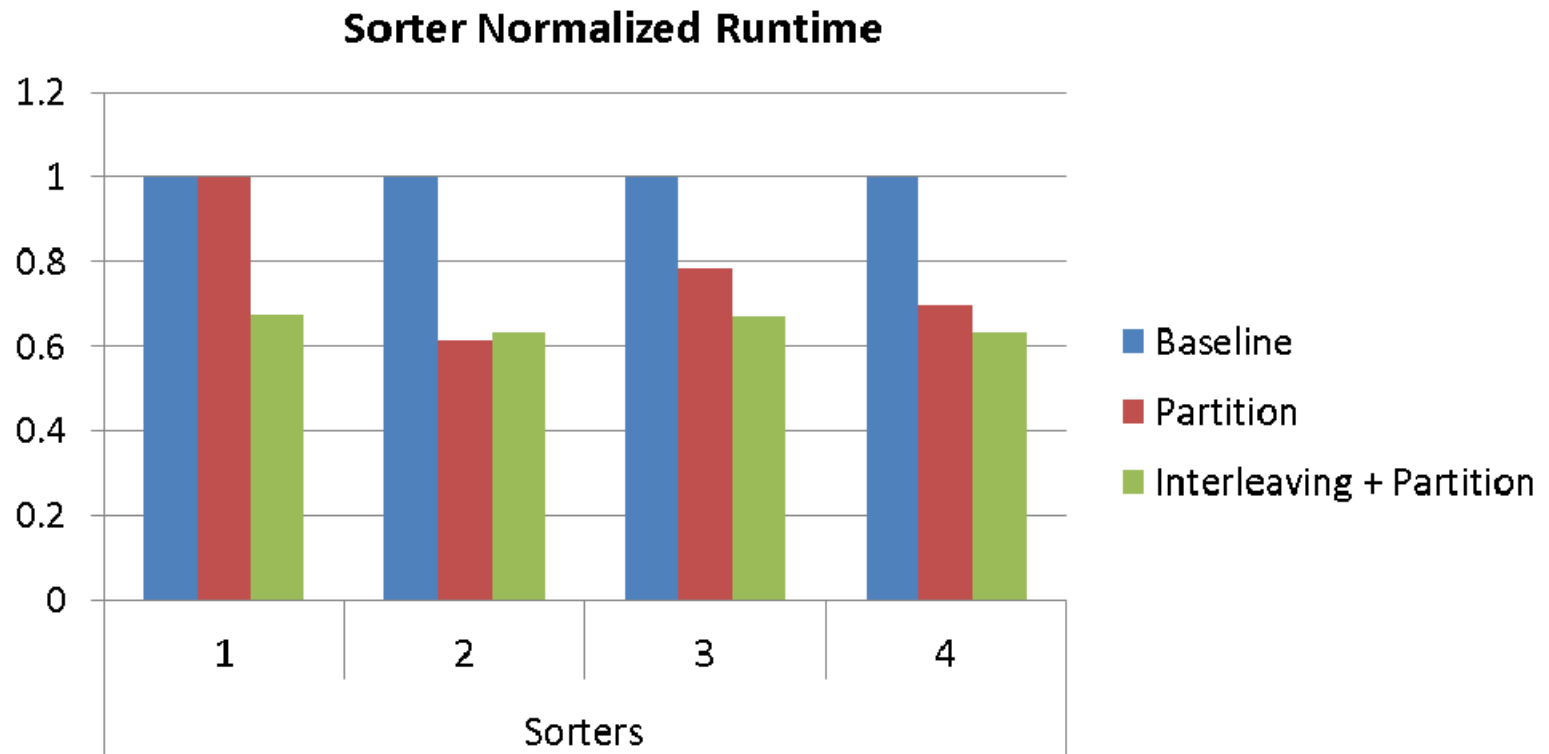
LEAP Memory Performance

- Memory interleaving



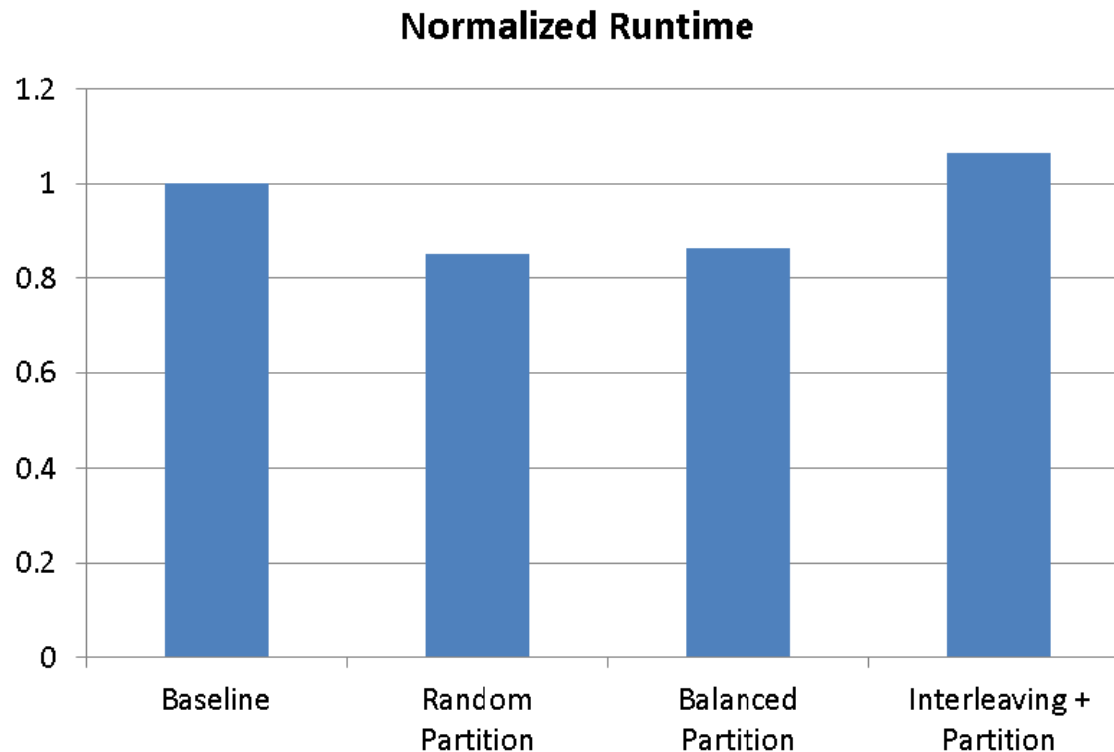
Case Study: Cryptosorter

- **Cryptosorter:** each sorter uses a LEAP private memory



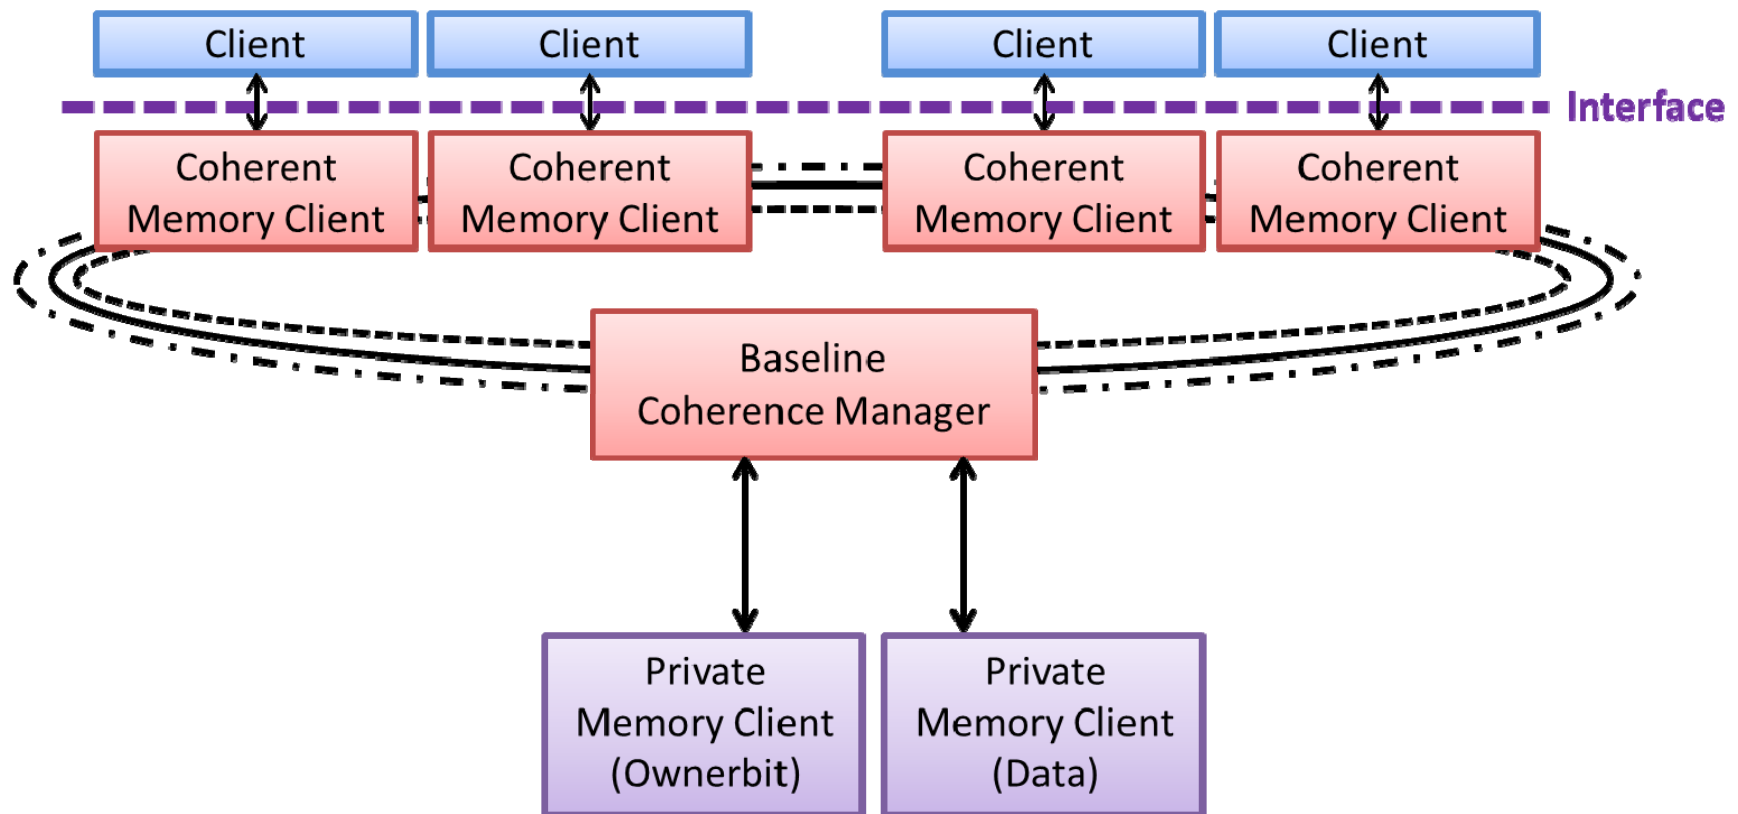
Case Study: Filtering Algorithm

- **Filtering algorithm for K-means clustering (HLS kernel)**
 - 8 partitions: each uses 3 LEAP private memories



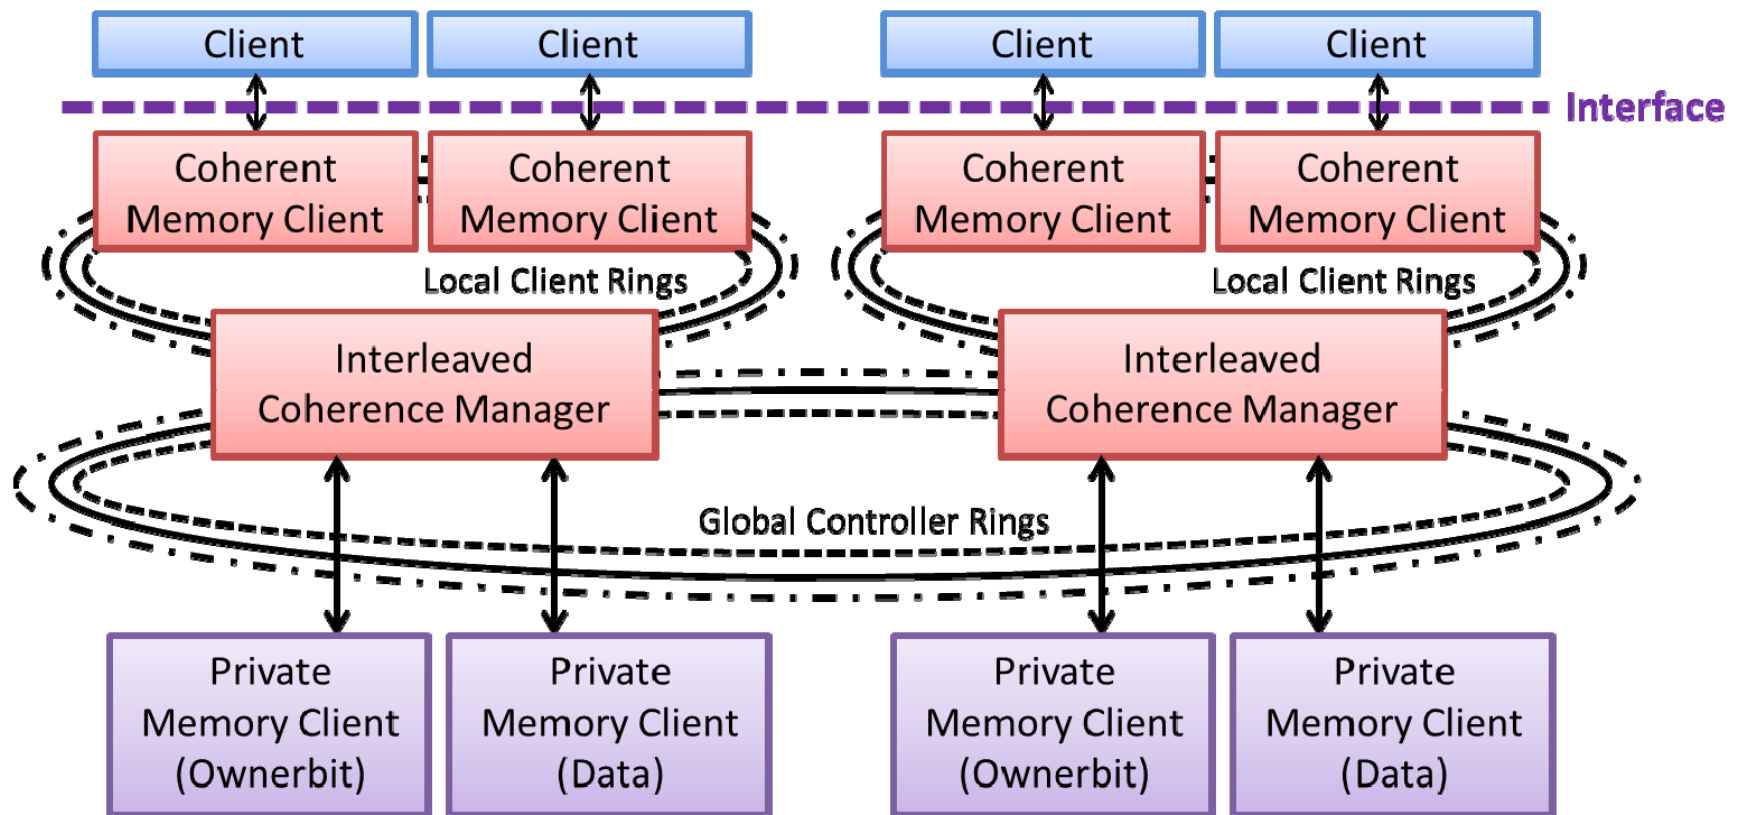
Coherent Cache Network Partitioning

- Baseline coherent memory



Coherent Cache Network Partitioning


- Coherent memory interleaving



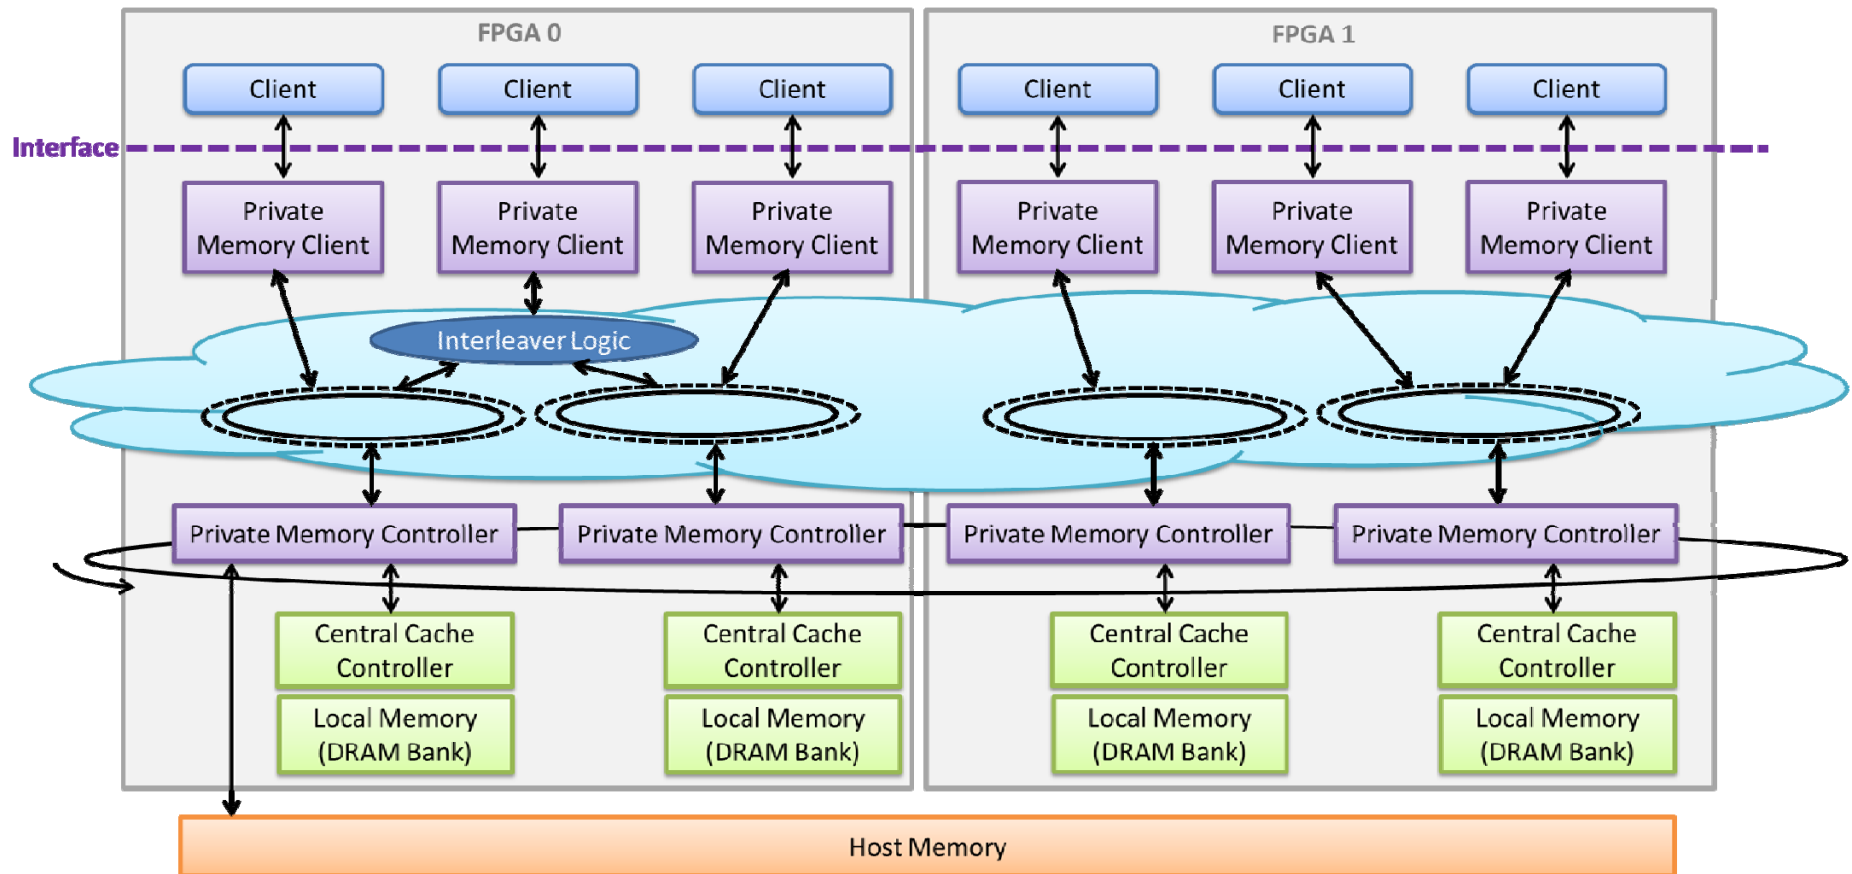
Private cache network optimization can be combined.

Case Study: Heat Transfer

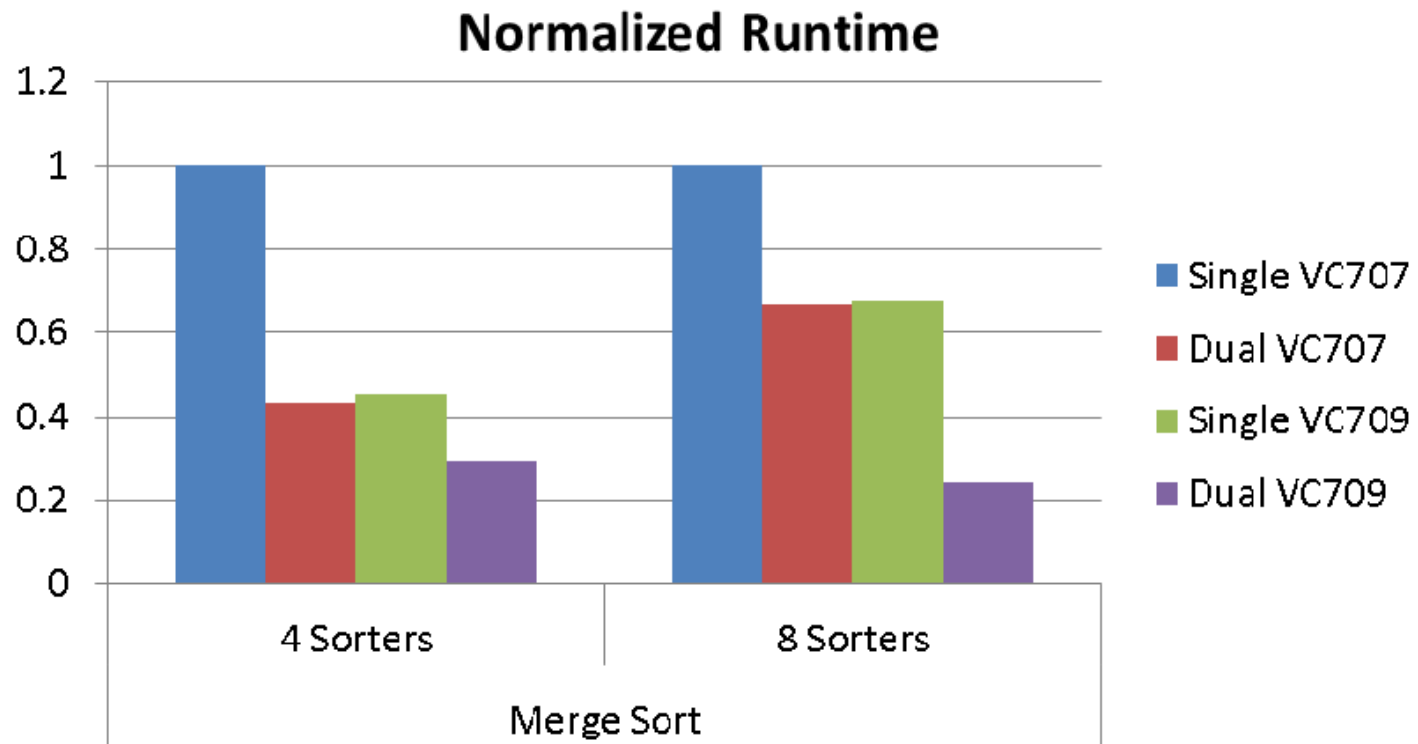
- Heat transfer: 16 engines, 1024x1024 frame



Moving to Multi-FPGA Platforms



Performance on Dual FPGAs



Conclusion

- We introduce the LEAP memory compiler that can transparently optimize the memory system for a given application.
- The compiler automatically partitions both private and coherent memory networks to efficiently utilize the increased DRAM bandwidth on modern FPGAs.
- Future work:
 - More case studies on asymmetric memory clients
 - More complex memory network topologies
 - Dynamic cache partitioning

Thank You